

Multilingual hierarchical classification of job advertisements for job vacancy statistics

Maciej Beręsewicz*, Marek Wydmuch†, Herman Cherniaev‡, Robert Pater§

Abstract

The goal of this study is to develop a multilingual classifier and conditional probability estimator of occupation codes for online job advertisements according in accordance with the International Standard Classification of Occupations (ISCO) extended with the Polish Classification of Occupations and Specializations (KZiS), which is analogous to the European Classification of Occupations. In this paper, we utilise a range of data sources, including a novel one, namely the Central Jobs Database, which is a register of all vacancies submitted to Public Employment Offices. Their staff members code the vacancies according to the ISCO and KZiS. A hierarchical multi-class classifier has been developed based on the transformer architecture. The classifier begins by encoding the jobs found in advertisements to the widest 1-digit occupational group, and then narrows the assignment to a 6-digit occupation code. We show that incorporation of the hierarchical structure of occupations improves prediction accuracy by 1-2 percentage points, particularly for the hand-coded online job advertisements. Finally, a bilingual (Polish and English) and multilingual (24 languages) model is developed based on data translated using closed and open-source software. The open-source software is provided for the benefit of the official statistics community, with a particular focus on international comparability.

Keywords: hierarchical multi-class classification; ISCO; ESCO; occupation; skills demand

*Poznań University of Economics and Business, Institute of Informatics and Quantitative Economics, Department of Statistics, Al. Niepodległości 10, 61-875 Poznań, Poland, E-mail: maciej.beresewicz@ue.poznan.pl; Statistical Office in Poznań, ul. Wojska Polskiego 27/29 60-624 Poznań, Poland.

†Poznan University of Technology, Institute of Computing Science, Piotrowo 2, 60-965, Poznań, Poland, E-mail: mwydmuch@cs.put.poznan.pl

‡Educational Research Institute in Warsaw, Poland

§University of Information Technology and Management in Rzeszów, and Educational Research Institute in Warsaw, Poland

1 Introduction

Online and administrative sources yield substantial data rich in pertinent information that can be leveraged to enhance the understanding of labour market dynamics, as demonstrated by previous research (Hersh and Harding 2018; Beręsewicz et al. 2021). In particular, data from online job advertisements (OJA) have been employed extensively in the tracking of changes in vacancies (see, for example, Beręsewicz and Pater 2021; Turrell et al. 2021; Hershbein and Kahn 2018). The analysis of such information allows for the examination of the impact of current global shocks on labour demand. While the estimation of the number of jobs is a crucial aspect, the analysis and forecasting of structural change represents a significant challenge in economics during labour market transformation. The analysis of skill demand necessitates a disaggregate approach, as the general definition of skills is unable to capture the specificities in firms' demand for skills.

Lightcast (<https://lightcast.io/about/data>), the largest global company providing statistics on OJA, classifies occupations but does not disclose the models or the quality of said models (e.g., classification errors). The European Union agency, Cedefop, is responsible for the collection of OJA data across all EU member states, which is then made available for analysis through the Skills-OVATE tool¹. Cedefop classification of occupations is based on the International Standard Classification of Occupations (ISCO) up to the unit group level. Initially, they employed Support Vector Machines (SVMs), which Cedefop identified as the optimal approach in terms of classification accuracy (cf. Boselli et al. 2017; Colombo, Mercorio, and Mezzanzanica 2018).

A more recent machine learning approach dedicated to job offers was proposed by Turrell et al. (2022). The researchers employed a corpus of 15.2 million advertisements posted between 2008 and 2016 on [Reed.co.uk](https://www.reed.co.uk) for the purpose of classifying a job offer according to the 3-digit codes of the UK Office for National Statistics' (ONS) Standard Occupational Classification (90 minor occupational groups). In order to achieve this, the researchers employed a fuzzy matching technique. The researchers report that manual assignment of occupations yielded 76% correctness (on a sample of 330 job offers), while their algorithm achieved 91% accuracy (on a sample of 67,900 job offers). However, the authors also indicate that their assessment was based on only 34% of the initially collected job offers due to the inability of the ONS algorithm to

produce a confident label. In their study, European Commission (2022) employed a pre-trained multilingual XLM-RoBERTa (Conneau et al. 2020) transformer model (a type of deep neural network; see, for example, Vaswani 2017) to map national occupational classifications to the European Skills, Competences, Qualifications and Occupations (ESCO) framework. In addition to national qualifications and the ESCO dictionary, the authors make use of descriptions drawn from the Qualification Dataset Register and European Employment Services (EURES) online job advertisements. The authors conclude that the model demonstrates promising results when applied to work in a variety of languages. Chen and Mason (2024) put forth the utilisation of explainable AI (XAI) methodologies to develop a lexicon-based approach. The method was applied to identify traditional and green energy jobs in online advertisements, with a reported precision rate of 82%. In 2024, Eurostat, in collaboration with Cedefop, launched a text classification competition at the ISCO level. They provided a hand-coded dataset comprising 26,000 OJA, with the highest score for the lowest common ancestor at 58% (for further details, see <https://statistics-awards.eu/>). It should be noted that neither the data nor the algorithms are disclosed, so they cannot be used by researchers or the official statistics community.

In this study, we propose a hierarchical conditional probability estimator and a classifier of job advertisements for the purpose of estimating job vacancy statistics. The methodology is then applied to predict the probability and classify the advertisements across occupations according to the KZiS, which organises occupations into a five-level hierarchy of job categories, represented by a six-digit code. Each digit encodes the child category, with the exception of the final level, which is encoded by two digits. The 1-digit code represents the widest major category, while the 6-digit code corresponds to the fifth level, the most narrow category. This imposes a classical hierarchical multi-class classification problem. Similarly to the approach taken in European Commission (2022), our classifier is based on a pre-trained transformer model, which we have carefully fine-tuned for the specific task of job classification. The approach was tested using XLM-RoBERTa and HerBERT, a transformed model pre-trained on a large multilingual corpus that included the Polish language (Mroczkowski et al. 2021). The classifier was trained and evaluated using a dataset of online job advertisements gathered from Polish online job boards and administrative records.

Our contribution can be summarised as follows:

- We propose and develop a hierarchical conditional probability estimator and classifier that takes into account the structure of the official classification of occupations. This approach may serve as a significant complement to traditional hand-coding techniques and facilitate the utilisation of big data sources for official statistics.
- We employ a novel data source: large-scale administrative data comprising job offers that have been hand-coded by experts from public employment offices. In addition to the data itself, we provide a detailed description of its quality, along with the experts' coding, which is often absent or unavailable in the existing literature.
- We developed a model for 24 languages included in the European Union, with a view to studying the performance of these languages.
- The open-source software and models are made available to the official statistics community so that they may benefit from the proposed methodology, either by utilising the classifiers or by modifying them to suit their own purposes.

The remaining part of the paper is as follows. Section 2 describes Polish Classification of Occupations and Specialities (hereinafter KZiS) and its relation to the International Standard Classification of Occupations (ISCO). Section 3 provides theory for the proposed approach that takes into account the hierarchical structure of official classifications. Section 4 describes the dataset used for training and testing, in particular administrative data sources, hand-coded online job offers and other available sources. Section 5 contains information about the quality of the data, in particular how accurate are the experts' coding. Section 6 presents experimental setup and its implementation. Section 7 presents results for the proposed classifiers for Polish, bilingual (Polish and English) and multilingual datasets. The paper concludes with a discussion and an exploration of the potential for extending the approach to other classifications, including the Occupational Information Network (O*NET), the European Skills, Competences, Qualifications and Occupations (ESCO) framework, and the Standard Occupational Classification (SOC). The supplementary materials provide a brief tutorial on the use of the software, as well as detailed results.

2 Classification of occupations and specializations

The KZiS classification is created by the Ministry of Family, Labour and Social Policy and provides basic information about job type ([Ministry of Family, Labour and Social Policy 2014](#)). It is based on ISCO established by the International Labour Organization (ILO). The current classification, available since 2023 is based on ISCO-08 and updated with new occupations, especially from the vocational education system, that are annually announced by the Ministry of National Education².

The KZiS is a 6-digit, hierarchically structured classification system. The hierarchy encompasses a range of occupational groups, from one-digit major groups to six-digit occupations (cf. Table 2.1). It develops the ISCO classification, which ends with 4-digit codes (unit groups). In this case, it is similar to the later established European Skills, Competences, Qualifications and Occupations (ESCO) classification. The KZiS classification encompasses over 2,500 occupations, which are grouped into ten major categories. In addition to the Armed Forces (group 0), the classification is primarily delineated by the level of job position (or skill) through a process of alignment with the UNESCO International Standard Classification of Education (ISCED 2013) and the levels of the Polish Qualifications Framework, as applicable to the majority of occupations within a given group.

Table 2.1: Structure of the classification of occupations and specializations (as of 01.01.2023)

(1 digit) Major Group	2 digits Sub-major Groups	3 digits Minor Groups	4 digits Unit Groups	6 digits	
				with the rest	without the rest
0 – Armed Forces Occupations	3	3	3	3	3
1 – Public Authorities, Senior Officials and Managers	4	11	31	202	172
2 – Professionals	6	31	99	789	705
3 – Technicians and Associate Professionals	5	20	87	610	535
4 – Clerical Support Workers	4	8	27	89	71
5 – Services and Sales Workers	4	13	39	166	137
6 – Skilled Agricultural, Forestry and Fishery Workers	3	9	17	63	54
7 – Craft and Related Trades Workers	5	14	69	476	412
8 – Plant and Machine Operators and Assemblers	3	14	41	387	349
9 – Elementary Occupations	6	11	32	126	111
All categories	43	134	445	2,911	2,549

Note: The term "the rest" is used to denote those occupations that are not classified elsewhere. To illustrate, the 4-digit occupation code "2522 Systems Administrators" is extended with two specific occupations: "252201 Computer Systems Administrator" and "252202 Integrated Management Systems Administrator". Additionally, the third category is "252290 Other Computer Systems Administrators".

The majority of occupations within major group 1 and 2 require individuals to have com-

pleted a tertiary education qualification. These roles are associated with the highest skill levels, which entail the worker engaging in intricate practical and technical tasks (skill level 3), problem-solving, decision-making, and creativity (skill level 4). The majority of occupations within Groups 3 to 5 typically necessitate upper secondary education, frequently in a vocational context. These occupations are associated with ISCO skill levels 2 and 3. Skill level 2 encompasses the ability to perform tasks such as operating machinery, mechanical and electronic equipment, vehicles, and the manipulation, ordering, and storage of information. In order to perform tasks associated with ISCO skill level 2, secondary education is required for occupations within groups 6 to 8. Positions within major group 9 typically require a secondary education at most. Workers in these roles typically engage with tasks at levels 1-3 of the ISCO skill spectrum. Level 1 is associated with the completion of routine and elementary physical and manual tasks. The following section presents an overview of our methodology, which incorporates a consideration of the hierarchical structure of the occupation classification.

3 Hierarchical multi-class classification

Hierarchical multi-class classification is a classical machine learning problem in which an object can be assigned a set of classes that are hierarchically organized in the form of a tree (Sun and Lim 2001; Silla and Freitas 2011), and a maximum of one class can be predicted per hierarchy level. Additionally, it is often required that every prediction must be coherent with respect to the hierarchy, meaning that for every predicted class, the classifier must also predict all its ancestors in the hierarchy. In other words, hierarchical multi-class classification can be seen as a problem of selecting *a single path from the tree root to a leaf*. Because we are interested in population analysis of the job market, we not only require the classifier to assign discreet classes but also provide probability estimates of an online job offer belonging to the specific class at each level of the hierarchy.

3.1 Formal setting

Let $\mathcal{C} = \{1, \dots, m\}$ be a finite set of m classes with hierarchical relations between them encoded using a prefix code. Any such code can be given in the form of a hierarchy tree in which a path from a tree root to a leaf node corresponds to a code word. Under the coding, each class $c \in \mathcal{C}$ is

uniquely represented by a tree node, that is identifiable by a code $\mathbf{z}^l(c) = (z_1(c), \dots, z_l(c)) \in \mathcal{Z}^l$, where l is the length of the prefix and \mathcal{Z}^l is a set of all possible codes of length l . Only the tree root has no class assigned to it, and its code is of length 0, $\mathbf{z}^0 = \emptyset$. The class c with the code $\mathbf{z}^l(c)$ of length l is an ancestor of a class c' with the code $\mathbf{z}^{l'}(c')$ of length $l' > l$ if and only if $\forall_{j=1}^l z_j(c) = z_j(c')$. If $l' = l + 1$, the class c is called a parent class of c' and c' a child class of c .

For $z_i \in \{0, 1\}$, the code and the hierarchy tree are binary. In general, the code alphabet can contain any number of symbols. Furthermore, z_i can take values from different sets of symbols depending on the previous values in the code. In other words, the code can result in nodes of a different parity even in the same tree. While this is not a case considered in this work, the tree does not need to be perfectly balanced, resulting in code words of different lengths for different leaf nodes. We present an example of a hierarchy tree in Figure 1.

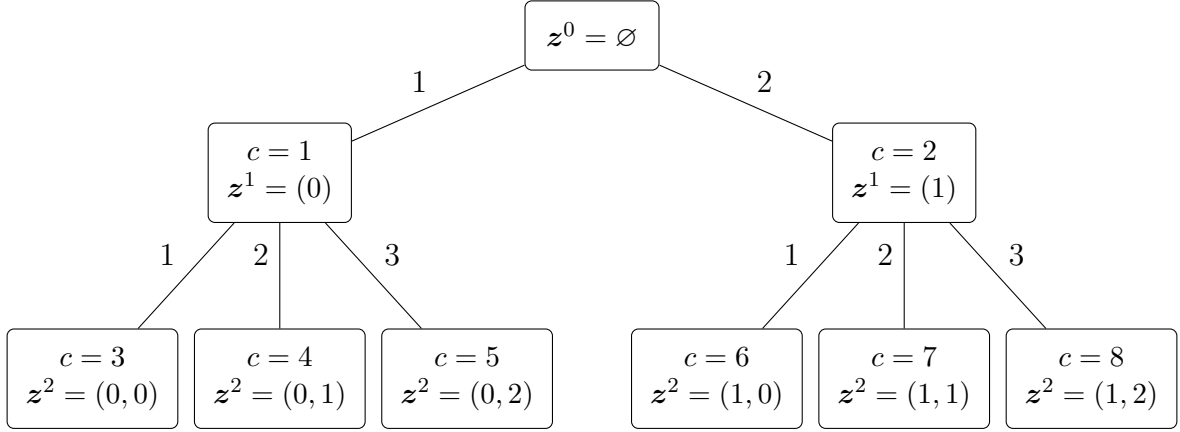


Figure 1: An example of class hierarchy with 8 classes and their assigned codes \mathbf{z} organized into two levels.

Let \mathcal{X} denote an instance space, and an instance $\mathbf{x} \in \mathcal{X}$ be associated with a single code word \mathbf{z} corresponding to a leaf node in the hierarchy tree, which corresponds to a subset of classes $\mathcal{C}_{\mathbf{x}}$ being assigned to the \mathbf{x} . This subset is often called a set of relevant labels, while the complement $\mathcal{L} \setminus \mathcal{L}_{\mathbf{x}}$ is considered as irrelevant for \mathbf{x} . We identify a set $\mathcal{L}_{\mathbf{x}}$ of relevant labels with a binary (sparse) vector $\mathbf{y} = (y_1, \dots, y_m)$, in which $y_c = 1 \Leftrightarrow c \in \mathcal{L}_{\mathbf{x}}$. By $\mathcal{Y} = \{0, 1\}^m$ we denote a set of all possible label vectors. We assume that observations (\mathbf{x}, \mathbf{y}) are generated independently and identically according to the probability distribution $\mathbb{P}(\mathbf{X} = \mathbf{x}, \mathbf{Y} = \mathbf{y})$ (denoted later by $\mathbb{P}(\mathbf{x}, \mathbf{y})$) defined on $\mathcal{X} \times \mathcal{Y}$.

Our goal is to find a *classifier* $\mathbf{h}(\mathbf{x}) = (h_1(\mathbf{x}), \dots, h_m(\mathbf{x}))$, which in general can be defined

as a mapping $\mathcal{X} \rightarrow \mathbb{R}^m$, that minimizes the *expected loss* (or *risk*):

$$L_\ell(\mathbf{h}) := \mathbb{E}_{(\mathbf{x}, \mathbf{y}) \sim \mathbb{P}(\mathbf{x}, \mathbf{y})} [\ell(\mathbf{y}, \mathbf{h}(\mathbf{x}))] , \quad (1)$$

where $\ell(\mathbf{y}, \hat{\mathbf{y}})$ is the (*task*) *loss*.

In particular, as our goal is to use the classifier for population analysis, we aim at estimating the conditional probabilities of classes $\eta_c(\mathbf{x}) := \mathbb{P}(y_c = 1 | \mathbf{x}) = \mathbb{P}(\mathbf{z}^l = \mathbf{z}^l(c) | \mathbf{x})$. To obtain the estimates of conditional probabilities, one can use the class-wise log loss as a surrogate objective:

$$\ell_{\log}(\mathbf{y}, \mathbf{h}(\mathbf{x})) = \sum_{c=1}^m \ell_{\log}(y_c, h_c(\mathbf{x})) = \sum_{c=1}^m (y_c \log(h_c(\mathbf{x})) + (1 - y_c) \log(1 - h_c(\mathbf{x}))) . \quad (2)$$

Then the expected class-wise log loss for a single \mathbf{x} (i.e., the so-called *conditional risk*) is:

$$\mathbb{E}_{\mathbf{y} \sim \mathbb{P}(\mathbf{y} | \mathbf{x})} [\ell_{\log}(\mathbf{y}, \mathbf{h}(\mathbf{x}))] = \sum_{c=1}^m \mathbb{E}_{\mathbf{y} \sim \mathbb{P}(\mathbf{y} | \mathbf{x})} [\ell_{\log}(y_c, h_c(\mathbf{x}))] = \sum_{c=1}^m L_{\log}(h_c(\mathbf{x}) | \mathbf{x}) . \quad (3)$$

Therefore, it is easy to see that the point-wise optimal prediction for the c -th class is given by:

$$h_c^*(\mathbf{x}) = \arg \min_h L_{\log}(h_c(\mathbf{x}) | \mathbf{x}) = \eta_c(\mathbf{x}) . \quad (4)$$

Let us notice that under class hierarchy and the assumption that only one leaf node is assigned to instance \mathbf{x} , the conditional probabilities of labels obey the following relations:

- The conditional probability of any parent class is equal to a sum of conditional probabilities of its children:

$$\mathbb{P}(\mathbf{z}^l | \mathbf{x}) = \sum_{z_{l+1}} \mathbb{P}(\mathbf{z}^l \cup (z_{l+1}) | \mathbf{x}) . \quad (5)$$

- Since only a single class per hierarchy level can be relevant, the sum of all conditional probabilities of classes on the same hierarchy level (with the same length of the codes l) should sum to 1:

$$\sum_{\mathbf{z}^l \in \mathcal{Z}^l} \mathbb{P}(\mathbf{z}^l | \mathbf{x}) = 1 . \quad (6)$$

3.2 Modeling probabilities under class hierarchy

We are not only interested in possible accurate estimates of conditional probabilities $\eta_c(\mathbf{x})$ for each class c , we additionally require for the predictions of the classifier to be coherent under class hierarchy, that is, obey the relations presented in Equations (5) and (6). Naively training the classifier using class-wise log loss (2) does not guarantee that class relations will hold. Because of that, in this work, we consider two classical approaches that guarantee coherency of prediction under class hierarchy.

We refer to the first approach as a “*bottom-up*” (cf. Barbedo and Lopes 2006). The bottom-up approach is one of the simplest to deal with the hierarchical classification problem. In this type of classifier, the conditional probability estimator of classes is trained only for classes in the leaf nodes of the hierarchy tree (i.e., 6-digit KZiS codes), and information about the class hierarchy is ignored. Then, during inference, after predicting probabilities for leaf labels, the predictions for labels higher in the hierarchy are reconstructed by summing the estimates of their direct children as in (5) until the root node is reached.

We refer to the second approach we consider as the “*top-down*” approach. In this algorithm, the probability of a given class is determined by a sequence of decisions made by node classifiers that predict subsequent values of the code word. By using the chain rule of probability, we obtain:

$$\eta_c(\mathbf{x}) = \mathbb{P}(\mathbf{z}^l(c)|\mathbf{x}) = \prod_{i=1}^l \mathbb{P}(z_i(c)|\mathbf{z}^{i-1}(c), \mathbf{x}). \quad (7)$$

The classifier then is formed by training probability estimator for every $\mathbb{P}(z_i(c)|\mathbf{z}^{i-1}(c), \mathbf{x})$. This can be done using the same log loss (2). By using logistic loss and a linear model $f_{\mathbf{z}^i}(\mathbf{x})$ in each node \mathbf{z}^i for estimating $\mathbf{P}(z_i | \mathbf{z}^{i-1}, \mathbf{x})$, we obtain the popular formulation of HSM. Let us notice that since we deal here with a multi-class distribution, we have that:

$$\sum_c \mathbb{P}(z_i = c | \mathbf{z}^{i-1}, \mathbf{x}) = 1 \quad (8)$$

This approach is known in the literature as the nested dichotomies (Fox 1997) in statistics, hierarchical softmax approach (Morin and Bengio 2005), conditional probability estimation trees (Beygelzimer et al. 2009) or probabilistic classifier trees (Dembczynski et al. 2016)

3.3 Statistical challenge of long-tail distribution

It is common in classification tasks with a large number of labels that labels are highly unbalanced – only a few classes have a large number of data points, while the vast majority is represented by a much smaller number of examples. This kind of distribution is said to be *long-tailed* (Bhatia et al. 2015; Babbar and Schölkopf 2017). The long-tail distribution imposes statistical challenges in modeling the conditional probability distribution using classical machine learning algorithms like logistic regression or naive Bayes classifier, as they are not performing well under such a low data regime.

In our task, the classifier must predict a set of classes out of 2,911 possible ones. Figure 2 illustrates the long-tail distribution present in the dataset, which is described in detail in Section 4. It is evident that almost one-third of the classes are represented by less than 10 examples, while only around 250 classes have more than 100.




Figure 2: Last level classes (codes) frequency in the study datasets. The X-axis shows the class rank when sorted by the frequency of positive instances, and the Y-axis shows the number of positive instances.

3.4 Proposed model architecture

To address long-tail distribution, instead of training the classifier from scratch, we opted for using a pre-trained transformer model. We are focusing on the popular family of models based on the Bidirectional Encoder Representations from Transformers (BERT; Devlin 2018)

architecture.

BERT is a pre-trained language model that has revolutionized natural language processing (NLP). It is an encoder-only model, meaning it is designed to create representations of an input text in the form of a vector that can be used for a variety of NLP tasks such as sentiment analysis, spam detection, or any other text classification task. BERT model uses bidirectional Transformer architecture. This mechanism helps the model weigh the importance of each word in the input sequence by considering its relationship with all other words. It allows BERT to capture context in a bidirectional way, meaning it considers both the left and right context when processing words.

It can be pre-trained on vast amounts of data in an unsupervised manner and can be fine-tuned for downstream tasks such as considered text classification. BERT models are usually pre-trained using two key strategies:

- Masked Language Modeling (MLM): A portion of the input tokens are randomly masked, and the model is trained to predict these masked tokens based on their context.
- Next Sentence Prediction (NSP): This task trains the model to understand sentence relationships by predicting whether two sentences appear consecutively in the original text.

In this work, we take a pre-trained BERT model and use its predicted text representation (hidden representation of [CLS] Token) combined with bottom-up or top-down output with linear models for modeling conditional probabilities of classes. We carefully fine-tune the resulting architecture in an end-to-end fashion – both partners of bottom-up and top-down outputs are updated, as well as the parameters of the BERT transformer model. This way, we obtain a model that is well-adjusted to the task while retaining its vast general language knowledge, that is, about words and concepts unseen in the training set used for fine-tuning. This should allow it to perform well also on a long-tail – classes with a small number of training examples. In the next section we describe the data used to train and test the proposed models.

4 Data sources

4.1 Official classifications and definitions

This dataset has been prepared for classification using the official classifications and accompanying documentation available on the website of the Ministry of the Family, Labour and Social Policy.

- Job description search engine³ – This service provides a comprehensive description of all occupations, including the name, code, synthesis, and a list of elementary and additional job tasks. For illustrative purposes, please refer to Figure 3. It is important to note that the descriptions of the occupations exhibit considerable variation in terms of both detail and completeness.
- INFOdoradca+ webservice⁴ – The service comprises a curated list of 1,000 occupations, offering a comprehensive overview of the profession or job role in question. It provides detailed information on the specific skills, qualifications and authorisations required for entry into the profession or job role, as well as the tasks and competencies necessary for success in the field. For an illustrative example, please refer to Figure 3, which depicts a description from the INFOdoradca+ web service.

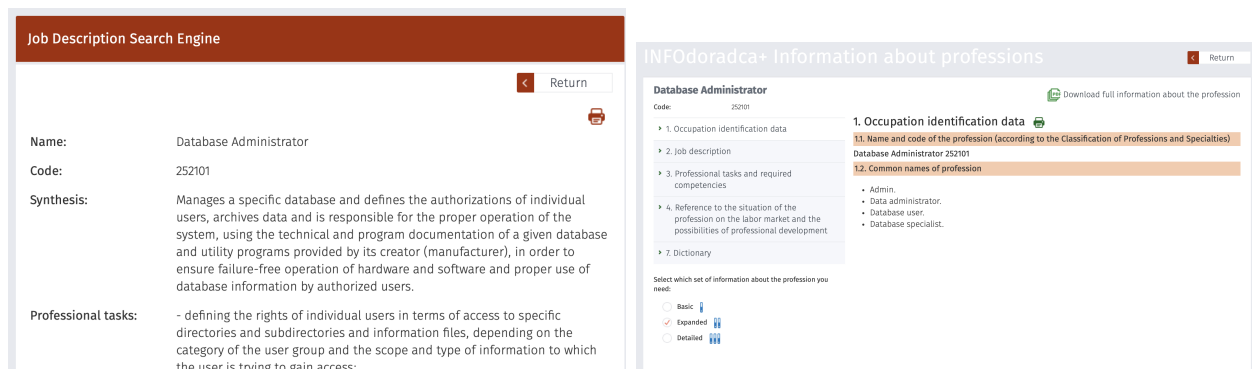


Figure 3: Example screenshots of the job description engine and INFOdoradca webservice

In addition to the official classifications, the following datasets were employed:

- thesaurus created by Statistics Poland⁵ – contains 1,338 unique occupations and their synonyms, with the number of synonyms varying from 2 (equivalent to approximately 40% of all occupations) to 82 (for one occupation 813101 "Aparatowy procesów chemicznych")

(Operator of chemical processes equipment)) with mean and median equal to 4.7 and 3 respectively.

- KZiS-ESCO linkage database – list of selected occupations from KZiS linked to ESCO proposed by Stęchły and Woźniak (2023). This dataset contains 557 unique codes along with its descriptions downloaded from the ESCO website⁶. In our dataset, we only used exact matches (about 8.7% of all KZiS codes).
- Civil service job offers database (KPRM) – 2,941 job offers from the civil service (<https://nabory.kprm.gov.pl>) for 12 occupation codes⁷ that meet the official definitions. This database was selected based on clerical review of matching job titles and descriptions.

For the official classifications (dictionary and INFOdoradca+), KZiS-ESCO and KPRM we have developed web-scraping algorithms that collected all available information on a given occupation.

In order to ensure the accuracy of online job advertisement classification, it is essential to utilise a comprehensive range of resources beyond official dictionaries or a limited number of advertisements from the Civil Service. Consequently, we have elected to employ the actual job postings from two primary sources: the Central Job Offers Database (CBOP), which encompasses all job advertisements submitted to Public Employment Offices (PEOs), and hand-coded job advertisements from a multitude of online sources. The subsequent subsections will delineate the characteristics of these datasets.

4.2 Central Job Offers Database

Public employment services in Poland include public employment offices (PEOs), which operate at NUTS4 level (*Nomenclature des unités territoriales statistiques, Nomenclature of Territorial Units for Statistics*) and are responsible for registering and managing unemployment. This means that there are two sources of data about employment in Poland: (i) the Labor Force Survey (LFS) and (ii) registered data collected by PEOs. The register maintained by PEOs contains detailed characteristics of unemployed persons and job offers. The number of registered unemployed may be higher than the LFS estimate because registration is connected with free health insurance and unemployment benefits.

The structure of job offers from PEOs differs from the general population of job vacancies. There is an over-representation of jobs from companies that have an incentive to advertise their vacancies through public employment offices, for example in the case of refunded internships or publicly-subsidised workplaces for the disabled (Beręsewicz, Cherniaiev, and Pater 2021). Public entities, in particular, are more willing to publish job offers in PEOs because they are often obliged to do so by their own internal regulations. Finally, low-paid jobs are more often sent to PEOs because people with lower qualifications often rely on public institutions to help them find a job. Better-paying jobs are more often advertised on job boards, in media or through private HR agencies, which charges fees for their services (Gałęcka-Burdziak and Pater 2015; Radzikowski 2023).

DETAILS

Back

Job offer: JUNIOR PRODUCTION EQUIPMENT OPERATOR
Number: StPr/24/0455
Added: Powiatowy Urząd Pracy w Złotowie

Print out

Export

Send via e-mail

Add to favorites

Recently viewed

Compare with the offer in Polish

TERMS OF WORK AND PAYMENT

Place of work:

Kujańska 10E, 77-400 Złotów, powiat: złotowski, woj: wielkopolskie

See on map

Responsibilities:

☐ supervision over the operation of the production line;
☐ arming aluminum guides;
☐ Wrapping guides;
☐ ongoing quality control;
☐ packing and appropriate labelling of materials transferred to the warehouse;
☐ maintaining order and tidiness at the workplace.

Type of contract:

Employment contract for a trial period

Working conditions:

Temporary work offered by agency:

NO

Working hours:

1

Shift:

three shifts

Number of working hours per week:

40

Number of hours working hours per month:

not applicable

Figure 4: Example job posting from CBOP

Job offers may be submitted via a paper or online form, following a highly structured format⁸. The requisite information encompasses the entity’s name, address, and REGON/NIP identifier; the job title; the number of vacancies; the type of contract; a description of the responsibilities and requirements for the applicants. An example of a job posting is provided in Figure 4. It is important to note that the option to specify the job occupation according to the official classification is available on the online form, though this is not a mandatory requirement. In accordance with the findings of our discourse with PEOs, it has been established that the majority of employers refrain from populating this field, and instead, PEOs personnel undertake the coding of job advertisements according to the official KZiS classification manually.

Regrettably, we are unable to present any numerical results to substantiate our claims, as this information is not accessible through CBOP database. In Section 5, we examine the quality of CBOP coding using a sample of job ads. Information about the occupation is policy-relevant. The Ministry of National Education annually announces the forecast of the demand for employees in vocational education occupations on the national and regional labour markets⁹. Based on the forecast the Ministry differentiates the educational part of the general subsidy for vocational education students between local government units. The regional governments consider the forecast in preparing an annual list of occupations for which the costs of vocational training for young employees is reimbursed. The forecast is also used by regional labour market councils in issuing opinions on the validity of education in occupations.

The CBOP data is accessible via the API for registered users¹⁰. Each job advertisement comprises 174 fields in JSON format, including an occupation code according to the KZiS and descriptions provided by the employer or recruitment agency. The data used in this study was downloaded on a daily basis from the beginning of 2022 to the end of 2023. The data was then subjected to a process of cleaning, whereby erroneous codes or descriptions with missing data in all fields were removed¹¹. Following this, the data set comprised over 822,000 ads with 2,468 unique occupation codes and job ads descriptions, based on the aforementioned criteria.

The distribution of KZiS codes was markedly uneven, with 145 codes having only a single example, 128 codes having two examples, and codes with over 5,000 examples¹². To reduce the size of CBOP database, a stratified sample of the CBOP dataset was created, with strata defined by KZiS code and the number of characters in the description. A total of 11,572 strata were formed. Table 4.1 presents information regarding the number of advertisements by the number of characters in the advertisement description.

Table 4.1: Number of CBOP ads by the number of characters in the description

Number of characters	Number of ads/examples
[0, 50]	6,628
(50, 100]	50,576
(100, 200]	196,240
(200, 300]	172,942
(300, 400]	118,598
(400, 500]	79,699
> 500	198,614

From each stratum, a simple random sample was selected, with the size determined as

$\max\{1, \lfloor 0.2 \times N_h \rfloor\}$, where N_h is the strata sample size and $\lfloor \cdot \rfloor$ denotes rounding to the nearest integer. The final sample size was 167,244 (around 20% of the CBOP dataset).

4.3 Hand coded datasets

4.3.1 The 10,000 ads dataset

For our study, we have used two datasets of hand-coded job advertisements. The first consisted of 10,000 job advertisements randomly selected from a large dataset of web-scraped advertisements. The second dataset consisted of 1,000 job advertisements selected from one of Poland's largest online recruitment services using specific keywords to match a range of occupations. We selected the IT sector for its diversity, particularly during the analysed period, and for jobs that are less prevalent in online recruitment portals. This study focuses on the 10,000 dataset.

As a basis for the sample, we have used 1,805,967 job ads scrapped between 2021-06-28 and 2021-11-08 from 18 sources¹³. This dataset contained information on the web service (the source), url, language of title and description according to a given website, job title, employer, region, job type, position level, contract type, salary, category (according to a given source), web-scraping date and job description. Note that some columns contained missing data (e.g., salary).

Before sampling, the data was cleaned to remove rows that had no job description (233,464 observations), had zero words or less than 5 characters¹⁴ (93 observations). Then we removed records with non-Polish descriptions. We have used Google's Compact Language Detector 3 available through the `clld3` package in R (Ooms 2024) to detect the language of the job description and identified 181,425 rows with non-Polish text. In the end, the final dataset used for sampling consisted of 1,390,985 records (77% of the initial dataset).

Next, we have created two stratification variables. The first was based on the identification of specific occupations: "sprzedawca" or "kasjer" (seller or cashier), "doradca klient" (customer consultant), "magazynier" (warehouse worker), "przedstawiciel" (representative), "pracownik produkcji" (production worker), "pomoc kuchenna" or "kucharz" (kitchen help or chef), "biuro" (bureau), "sprzątac" or "sprzatacz" (cleaner with and without Polish-specific letter) and "specjalista" (specialist). If none was matching "other" category was assigned.

The second stratification variable was based on the number of words in the description

that were categorized into the following groups: $[1, 10]$, $(10, 25]$, $(25, 50]$, $(50, 100]$, $(100, 150]$, $(150, 200]$, $(200, 300]$ and $300+$. Table A.2 presents information about stratification and population sizes within these strata. The largest strata was "specjaliści" (specialists) with over 130k ads. The second largest were "sprzedawcy" (sellers) with close to 55k ads. The smallest group were "sprzątacze" (cleaners) with over 7k ads. According to the second stratification variable groups with over 50 words were more or less of the same size. There were about 40k ads with at most 10 words and 70k with at most 25 words.

In order to code the dataset, three experts from PEOs were recruited in February 2022. Due to budgetary limitations, the sample size was set at 10,000, with each expert responsible for coding 3,333 advertisements. However, due to rounding, the final sample size was 10,002. The sample size was selected using a stratified systematic sampling design, with probabilities proportional to the number of characters in a given ad. The sample was selected using `strata` from the `sampling` package (Tillé and Matei 2023). Further details regarding the sample size according to the source can be found in Table A.3. The primary source was `pracuj.pl`, `praca.egospodarka` and `aplikuj.pl`. It should be noted that 241 ads were from CBOP, which enabled an assessment of the quality of CBOP encoding.

A total of 3,233 advertisements were selected for each expert, with an additional 100 advertisements sampled for cross-validation of expert coding¹⁵. Each expert was given an Excel spreadsheet, available via Google Sheets, containing all the relevant information. The experts were tasked with coding the advertisements into six-digit codes and providing comments where necessary. They were given a one-month period to complete the coding of the advertisements into occupation codes.

4.3.2 The 1,000 ads dataset

In January 2023, an additional dataset of 1,000 advertisements was selected for manual coding by a single expert. The sample was selected from a comprehensive dataset of over 543,000 ads scraped from one of the largest services in Poland. In the case of this particular dataset, we were able to identify in excess of 227,000 unique descriptions. We then proceeded to implement a sample selection process in accordance with the following scheme: the specification of a search phrase (including regular expression notation) and the subsequent selection of the most

appropriate samples. The information pertaining to the aforementioned phrases, the selected advertisements and the number of cases in one of the largest services in Poland is detailed in Table A.4.

The rationale behind the selection of the phases is as follows. Firstly, the decision was taken to include advertisements for positions related to computer science, data science and information technology. Secondly, the intention was to incorporate occupations that were not previously covered by our existing sources, such as pharmacist, sworn translator or piano repairer.

4.4 The multilingual dataset

All datasets (over 200k rows) were translated into English and then to 22 other languages: Bulgarian, Czech, Danish, German, Greek, Spanish, Estonian, Finnish, French, Irish, Hungarian, Italian, Lithuanian, Latvian, Dutch, Portuguese, Romanian, Russian, Slovak, Slovenian, Swedish, Ukrainian. This procedure was done in two steps.

Firstly, the Polish text was translated into English using Google Translate. This was achieved by utilising the Google Sheets application and the `googletranslate()` function. This approach was employed for the translation of over 200,000 examples. Additionally, Argos Translate was employed for the translation of the remaining text. The open-source offline translation library, written in Python, is available at <https://github.com/argosopentech/argos-translate>. It uses the `OpenNMT` library (Klein et al. 2017), which in turn uses the `CTranslate2` library, which is a fast inference engine for Transformer models. The objective was to evaluate the influence of an automated language library on the classification of a bilingual dataset.

Secondly, the ads translated to English with Google Translate descriptions were translated into 22 languages using the Argos Translate software. The principal rationale for this approach was the scarcity of human resources (there were no assistants available to undertake the translation process using Google Sheets) and time constraints, given that access to GPUs was available, as supported by the Argos Translate open-source library.

Due to budgetary constraints, an assessment of the quality of automated translation was not conducted. However, based on a discussion with PEO experts, it was evident that if they were to code English or other language-based descriptions, they would likely utilise Google Translate or other free translation services.

5 Quality of the input data

5.1 Hand-coded data quality

5.1.1 The 1,000 dataset

The 1,000 dataset was coded by a single expert, and thus, cross-validation was not available. The expert identified 968 advertisements for which a single occupation code was deemed appropriate, while 32 advertisements were assigned multiple codes (27 with two codes, three with three codes, and one with four codes). Therefore, the final dataset comprised 1,035 advertisements with 226 codes, with the most prevalent being: "Programista aplikacji" (Applications programmer, 125 occurrences; 251401), "Analityk baz danych (data scientist)" (Data scientist, 77 occurrences; 252102) and "Specjalista do spraw cyberbezpieczeństwa" (Cybersecurity specialist, 63 occurrences; 252902). See Table A.5 for the information on the number of unique 6-digit and 1-digit codes for a given stratum defined by the search phrases.

5.1.2 The 10,000 dataset

The process of assessing the quality of the hand-coded data was an iterative one. Information about the coding is presented in Table 5.1. From the final sample, 32 advertisements with a description in a language other than Polish were removed, although they had been coded by experts.

Table 5.1: Information about the quality of hand-coded data

Initial sample	Non-Polish description	Non-matching codes	Multiple codes	Final sample
10,002	32	38	59	9,932

It was observed by experts that for 38 advertisements (0.38%; 13 from the first expert, 13 from the second, and 2 from the third), it was not possible to ascertain the occupation due to a lack of pertinent information or because no occupation in the official classification matched the description. In one instance, the expert indicated that the advertisement describes two distinct occupations. In 59 cases (representing 0.6% of those with codes), experts noted that multiple codes align with the description. Rather than removing these 59 cases, we elected to flag them as potential codes for the same description. Consequently, the final dataset comprised

9,932 ads (prior to accounting for multiple codes).

The subsequent stage of the study involved an investigation into the quality of the coding produced by experts based on a set of 100 ads that exhibited a certain degree of overlap. Two of the experts did not provide a code for one of the advertisements, whereas the others did. Accordingly, for the purposes of comparison, the total number of ads considered was 98. Table 5.2 presents the data on the rate of agreement of expert coding at the 1-, 2-, 4-, and 6-digit levels (section "before clerical review") along with 95% confidence interval using sampling weights using the `survey` (Lumley 2004) package. In addition, we have estimated Cohen's kappa coefficient for agreement between 1 digit codes for each pair of experts (we used `survey::svykappa()` function).

Table 5.2: Point and 95% interval estimates of the rate of agreement of expert coding at 1, 2, 4, and 6 digits

Experts	1 Digit	2 Digits	4 Digits	6 Digits	Kappa
Before clerical review					
1 & 2	87.7 (80.5, 93.4)	84.2 (76.6, 90.5)	78.7 (70.3, 86.1)	74.2 (65.2, 82.2)	85.4 (77.8, 93.0)
1 & 3	86.2 (78.0, 92.7)	86.2 (78.0, 92.7)	74.7 (65.2, 83.2)	66.2 (56.1, 75.7)	83.5 (75.1, 92.0)
2 & 3	85.9 (77.7, 92.5)	84.5 (76.1, 91.3)	75.6 (66.4, 83.8)	68.5 (58.7, 77.6)	83.1 (74.4, 91.8)
All	79.9 (70.9, 87.6)	78.5 (69.4, 86.3)	66.1 (56.3, 75.3)	59.2 (49.1, 69.0)	—
After clerical review					
1 & 2	92.6 (86.2, 97.1)	90.3 (83.4, 95.5)	84.4 (76.4, 91.0)	80.5 (72.0, 87.8)	91.2 (84.8, 97.6)
1 & 3	90.2 (83.1, 95.5)	90.2 (83.1, 95.5)	84.7 (76.4, 91.5)	77.5 (68.2, 85.6)	88.3 (81.1, 95.6)
2 & 3	86.8 (79.0, 93.0)	85.8 (77.8, 92.2)	80.8 (72.0, 88.3)	72.9 (63.3, 81.5)	84.2 (76.1, 92.4)
All	84.8 (76.6, 91.5)	83.8 (75.4, 90.7)	75.6 (66.2, 83.9)	68.5 (58.8, 77.5)	—

In general, pairs of experts agree at the level of 86% at the first digit of occupation codes, with this figure decreasing to 73% and approximately 67% at six digits. The level of agreement among all experts is 80% for the first digit and decreases to approximately 60% at the sixth digit. The findings were presented to the experts, who were invited to provide comments and suggest potential amendments. Expert 1 elected to alter 10 of their original codes, Expert 2 chose to modify 13 of their initial codes, and Expert 3 revised six codes (including the addition of two missing codes). The results of the clerical review are presented in Table 2 (section 'After clerical review'). As anticipated, the agreement rate increased, although some discrepancies between experts remained. The overall agreement on a single digit increased to 85%, while for six digits it reached 68%. The primary source of discrepancy was the lack of relevant information about the education or skills. For all pairs of experts, the Kappa coefficient for agreement of 1 digit codes shows strong agreement with levels close to 85%.

Next, we focus on the assessment of the quality of CBOP coding, as in our sample, we had 236 ads from this source. We were able to link 213 unique ads to the CBOP using the hash identifier present in the ad’s URL. Table 5.3 presents information on the quality of CBOP coding according to the experts (point and 95% CI for the rate of agreement) and Kappa coefficient for each expert and CBOP as well as all experts (treated as one group) and CBOP as another (the last row denoted as *All*). Note that due to sparse data, all measures were calculated without the 1 and 6 occupation group for both CBOP and experts (3 obs were dropped).

Table 5.3: Point and 95%interval estimates of the rate of agreement of expert coding at 1, 2, 4 and 6 digits for CBOP ads

Expert	Count	1 Digit	2 Digits	4 Digits	6 Digits	Kappa
1	75	78.2 (65.9, 88.3)	74.9 (62.9, 85.2)	63.2 (50.7, 74.9)	56.3 (43.7, 68.5)	73.0 (59.2, 86.8)
2	66	79.0 (67.5, 88.5)	76.4 (64.6, 86.3)	72.2 (60.3, 82.7)	65.6 (53.5, 76.7)	74.3 (61.8, 86.8)
3	69	81.0 (70.3, 89.7)	80.0 (69.3, 89.0)	71.6 (59.8, 82.1)	62.7 (50.1, 74.5)	77.6 (66.2, 88.9)
All	210	79.3 (73.0, 85.0)	77.0 (70.6, 82.8)	68.8 (62.0, 75.2)	61.4 (54.4, 68.1)	75.2 (67.9, 82.4)

Subsequently, an assessment of discrepancies in coding was conducted by experts. Two experts consented to participate in the study and provided comments. The second expert determined that six of the 29 advertisements provided by CBOP were accurate, differing from the expert’s initial assessment. Additionally, four of the 29 advertisements exhibited a dual accuracy, with both codes being correct. Consequently, the expert’s agreement rate at the 6-digit level would increase to 76%, while at the 1-digit level, it would increase slightly, given that the disagreements are at more detailed levels. The third expert determined that in seven of the 30 advertisements, both codes are accurate. However, the expert highlighted that the lack of sufficient information precludes the ability to distinguish between the two. Consequently, for this expert, the agreement rate at the six-digit level would increase to 72%, while at the one-digit level, a slight increase is observed due to the disagreements occurring at more detailed levels. This also results in lower Kappa coefficients around 75% in comparison to about 85% presented in Table 5.2 but they are not significantly different if we compare confidence intervals.

Table 5.4 presents a comparison of agreements at the one-digit level. The figures in the table refer to the precision of the classification of CBOP ads. With the exception of the 4th group, the agreement rate for groups 2 to 9 is above 70%. However, as there is only one example in the first group, we refrain from commenting on it. Overall, it can be observed that the CBOP codes are characterised by a high level of confidence, with the main discrepancies occurring

between the second and fourth groups.

Table 5.4: Comparison of codes at the 1-digit groups between CBOP and experts coding

CBOP	Hand-coded									N
	1	2	3	4	5	6	7	8	9	
1	0.0	0.0	100.0	0.0	0.0	0.0	0.0	0.0	0.0	1
2	0.0	80.8	12.4	0.0	6.8	0.0	0.0	0.0	0.0	19
3	0.0	23.0	65.8	7.0	0.0	0.0	4.2	0.0	0.0	20
4	0.0	21.8	3.9	68.5	0.0	0.0	0.0	0.0	5.9	18
5	0.0	0.0	2.1	4.6	87.1	0.0	0.0	0.0	6.2	32
6	0.0	0.0	0.0	0.0	0.0	100.0	0.0	0.0	0.0	1
7	0.0	0.0	6.7	0.0	0.0	0.0	84.9	2.1	6.3	48
8	0.0	0.0	0.0	0.0	0.0	0.0	12.3	83.9	3.8	35
9	1.7	0.0	0.0	0.0	0.0	0.0	23.7	1.5	73.1	39

The assessment of the quality of hand-coded data is of paramount importance in the construction of a classifier, as the input codes (labels) must be accurate. In our approach, we utilise the codes in their original form; however, it is essential to consider the potential for uncertainty in the preparation of training datasets when interpreting the final results. Further research and the development of algorithms that take this into account will be required to address this issue in greater depth.

5.2 General description of the data

In this section, we present information on the quality of input data. The combined dataset contained 200,875 cases out of which 9,200 cases refer to the official dictionary¹⁶. Table 5.5 presents the coverage of occupation 6-digits occupation codes by main groups and source. The last row presents information about the number of rows for a given dataset. Datasets varied in terms of coverage. The most complete was CBOP with coverage close to 80% of all 6-digit codes within given occupation group. Due to sample size, hand-coded data covered only a small fraction of available codes. The combined dataset was missing 186 occupation codes¹⁷. It should be noted that for the training dataset the complete dictionary with all 2,911 codes was used.

Another factor that should be taken into account is the varying number of cases for each code. For instance, the largest number of cases was present for Industrial and craft workers (7) and Professionals (2) with close to 40k and 34k cases respectively. The less represented were Public officials, senior officials and managers (1), Farmers, horticulturists, foresters and fishermen (6), and Armed Forces (0).

Table 5.5: Coverage of 6 digit codes by main groups and sources (excluding official dictionary)

Group	Info+	GUS	ESCO	KPRM	Hand (1)	Hand (2)	CBOP	All	Cases
0	0.0	100.0	33.3	0.0	0.0	0.0	100.0	100.0	28
1	6.4	85.6	19.8	1.0	42.1	15.3	79.7	95.5	5,369
2	23.8	26.0	20.7	0.6	39.3	14.4	77.1	89.4	33,441
3	35.7	48.0	20.0	0.7	29.7	4.8	83.6	92.5	22,229
4	53.9	21.3	20.2	1.1	53.9	9.0	95.5	98.9	18,710
5	61.4	28.9	27.1	0.0	38.0	6.6	89.2	97.0	29,167
6	69.8	57.1	12.7	0.0	11.1	0.0	85.7	98.4	1,769
7	46.2	46.6	16.8	0.0	39.3	4.8	89.5	94.3	39,820
8	40.8	82.4	15.2	0.0	25.3	2.1	91.0	97.9	21,761
9	4.0	15.9	16.7	0.0	47.6	1.6	96.0	96.0	28,581
N	7,007	1,342	2,114	2,941	9,992	1,035	167,244		200,875

Source: Own elaboration. Group (1 digit) as in Table 2.1.

Table 5.6: Number of words for a given group for the combined dataset (official dictionary included)

Group	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	N
0	4.0	13.2	36.5	52.8	54.0	208.0	28
1	1.0	35.0	83.0	131.7	183.0	2,561.0	5,369
2	1.0	30.0	59.0	107.0	141.0	2,836.0	33,441
3	1.0	29.0	56.0	96.6	118.0	2,792.0	22,229
4	1.0	25.0	45.0	80.4	83.0	3,385.0	18,710
5	1.0	20.0	32.0	50.2	55.0	3,215.0	29,167
6	1.0	18.0	34.0	75.7	67.0	1,229.0	1,769
7	1.0	19.0	32.0	52.6	57.0	1,924.0	39,820
8	1.0	22.0	36.0	56.2	62.0	1,459.0	21,761
9	1.0	19.0	31.0	42.3	52.0	1,686.0	28,581
Overall	1	22	39	70	75	3,385	200,875

Source: Own elaboration. Groups as in Table 2.1.

Sources varied in terms of a number of words describing a given code. Table 5.6 presents information on the distribution of the number of words in the combined dataset for each group separately. The median number of words was 39 and varied between from 31 for the 9th to 83 for the 1st group. A comparison of the number of words in the descriptions from the two sources reveals that the CBOP descriptions contain, on average, approximately 50 words, while the descriptions from online sources contain more than 120 words on average. The detailed information for each input dataset is presented in the Supplementary Material in Table Table A.6.

Finally, for building the classifier, we have created training and test data according to the

following algorithm:

Step 1: Combine datasets: INFO+, KPRM, HAND (10k), HAND (1k), ESCO and CBOP into one dataset.

Step 2: item Calculate the size of 70% of cases stratified by source and 6-digit code.

item[Step 3:] Strata with only a single case were removed and used for testing.

item[Step 4:] Strata with > 1 cases were sampled using a stratified design with sample sizes defined as in Step 2; 70% for training and 30% for testing.

item[Step 5:] The final training dataset was created by combining 70% of the dataset from Step 4 with an official dictionary and GUS thesaurus.

item[Step 6:] The final test dataset was created by combining the dataset from Step 3 with 30% sample from Step 4.

Table 5.7 presents information on the number of cases and 6-digit codes in train and test datasets by source. The training dataset contains close to 143k cases while testing 58k. The main contribution is from the CBOP dataset. The number of codes in the training dataset is equal to the number of codes in the official classification, while in the test data not all codes are available.

Table 5.7: Number of cases and 6-digit codes in train and test data by source

Source	Train			Test		
	Cases PL	Cases ML	# Codes	Cases PL	Cases ML	# Codes
Official	9,200	22,143	2,911	-	-	-
Thesaurus	1,342	2,687	1,338	-	-	-
CBOP	116,879	280,461	2,213	50,365	181,648	2,468
ESCO	1,531	3,642	557	583	2,318	557
Hand 10k	6,720	16,081	708	3,272	12,050	1,039
Hand 1k	632	1,504	115	403	1,523	226
Info	5,004	12,010	996	2,003	7,012	996
KPRM	2,058	4,926	12	883	3,170	12
All	143,366	343,454	2,911	57,509	207,721	2,625

Note: Cases PL refers to Polish dataset, Cases ML refers to multilingual (24 languages) dataset

In order to verify the approach we translated train and test data from English (based on Google Translate) to 22 languages (all UE languages) with Argos Translate. This dataset over 3,4 mln records for training and close to 1,4 mln for test data. In order to avoid overfitting of the multilingual classifier we selected a stratified random sample of both datasets independently

with 10% and 15% from each stratum created by 6-digit occupation code and source type. Information about the number of cases for each dataset (train and test) are presented in Table 5.7 while detailed information about the distribution of languages are presented in Table 5.8.

Table 5.8: Distribution of languages in train and test data (in thousands of records) in the multilingual dataset and information about the number of tokens (in millions) on which XLM-Roberta was trained

Languages A-I				Languages I-U			
Language	Train	Test	XLM-RoBERTa	Language	Train	Test	XLM-RoBERTa
Bulgarian (bg)	12.3 (3.9%)	7.7 (4.0%)	5.5 (4.4%)	Italian (it)	12.2 (3.9%)	7.6 (4.0%)	5.0 (4.0%)
Czech (cs)	13.3 (4.2%)	11.1 (5.8%)	2.5 (2.0%)	Lithuanian (lt)	13.8 (4.3%)	7.2 (3.8%)	1.8 (1.5%)
Danish (da)	15.4 (4.9%)	7.4 (3.9%)	7.8 (6.2%)	Latvian (lv)	14.4 (4.6%)	8.1 (4.2%)	1.2 (1.0%)
German (de)	15.7 (4.9%)	8.0 (4.2%)	10.3 (8.2%)	Dutch (nl)	12.9 (4.1%)	7.6 (4.0%)	5.0 (4.0%)
Greek (el)	14.5 (4.6%)	8.3 (4.3%)	4.3 (3.4%)	Polish (pl)	12.2 (3.9%)	7.3 (3.8%)	6.5 (5.2%)
English (en)	14.2 (4.5%)	9.3 (4.8%)	55.6 (44.3%)	Portuguese (pt)	14.4 (4.5%)	10.2 (5.3%)	8.4 (6.7%)
Spanish (es)	16.0 (5.1%)	7.9 (4.1%)	9.4 (7.5%)	Romanian (ro)	12.3 (3.9%)	9.8 (5.1%)	10.4 (8.3%)
Estonian (et)	16.5 (5.2%)	7.7 (4.0%)	0.8 (0.7%)	Russian (ru)	12.5 (4.0%)	9.5 (5.0%)	23.4 (18.7%)
Finnish (fi)	14.6 (4.6%)	7.9 (4.1%)	6.7 (5.4%)	Slovak (sk)	16.4 (5.2%)	10.5 (5.5%)	3.5 (2.8%)
French (fr)	18.8 (5.9%)	11.3 (5.9%)	9.8 (7.8%)	Slovenian (sl)	13.4 (4.2%)	8.2 (4.3%)	1.7 (1.3%)
Irish (ga)	14.2 (4.5%)	8.0 (4.2%)	0.1 (0.1%)	Swedish (sv)	14.1 (4.4%)	10.2 (5.4%)	0.1 (0.1%)
Hungarian (hu)	15.0 (4.7%)	9.3 (4.9%)	7.8 (6.2%)	Ukrainian (uk)	14.4 (4.5%)	7.8 (4.1%)	0.0 (0.0%)

The number of cases for a given language varies between the two sources. To illustrate, in the case of hand-coded test data, Polish was the least represented language, with 287 cases, while French was the most represented, with 1,348.

Table 5.8 contains information about the number of tokens used to train XLM-RoBERTa. The primary objective is to demonstrate that some languages are markedly under-represented in the XLM-RoBERTa dataset, including Estonian, Irish, Swedish and Ukrainian, while others are over-represented, such as German and Russian. This may suggest that the quality of classification may vary between languages.

6 Experimental setup

6.1 Models and implementation

In the empirical study, we considered the following model variants of bottom-up and top-down approaches described in Section 3.2:

- As a baseline, we use bottom-up and top-down with simple linear models (i.e. logistic regression with linear part) for modelling probabilities using TF-IDF representations of texts.

- We compare it with the transformer-based architecture as described in Section 3.4, in both bottom-up and top-down variants fine-tuning them based on the following pre-trained models:
 - HerBERT-base (<https://huggingface.co/allegro/herbert-base-cased>, with 110M parameters) and HerBERT-large (<https://huggingface.co/allegro/herbert-large-cased>, with 336M parameters) – fine-tuned BERT for Polish language developed by Allegro,
 - XLM-RoBERTa-base (<https://huggingface.co/FacebookAI/XLM-roberta-base>, with 279M parameters) and XLM-RoBERTa-large (<https://huggingface.co/FacebookAI/XLM-roberta-large>, with 561M parameters) – multilingual models developed by Facebook/Meta.

As the baseline implementation, we use `napkinXC` (Jasinska-Kobus et al. 2020), which is an extremely simple and fast library for extreme multi-class and multi-label classification that implements, among the others, linear one-versus-rest multi-class classifier (used for bottom-up model) and hierarchical softmax (for top-down). In order to facilitate the implementation of our approach, we have developed an open-source software tool that enables users to either utilise existing models or train new ones.

For the implementation of the Transformer architecture, we use PyTorch (Paszke et al. 2019), PyTorch Lightning (Falcon and Team 2019), and Huggingface Transformers (Wolf 2019) libraries. We fine-tune transformer-based models using AdamW optimizer (Loshchilov and Hutter 2017), with the base learning rate equal to $1e-5$, weight decay to 0.01 (for all layers except bias terms and layer norms layers), and Adam hyperparameters to $\beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 1e-8$. We use a linear learning rate schedule with a warm-up phase of 50 training steps. The training was done on a machine with two Nvidia RTX5000 GPUs with 24 GB RAM each. For base models, we used 10 epochs, while for large, only 5. The reason for that was the computational time, for instance, for `HerBERT-base` one epoch was about 30 min, for `XLM-RoBERTa-base` about an hour, for `HerBERT-large` about 1.5 hours and `XLM-RoBERTa-large` about 2.5 hours. Training for the top-down approach was approximately two times longer than for the bottom-up approach.

6.2 Evaluation metrics

To evaluate the quality of obtained classifiers, we primarily report a per hierarchy level logistic loss (cross-entropy) to investigate if the proposed approach improves the quality of obtained estimates of conditional probabilities of classes $\eta_c(\mathbf{x}) := \mathbb{P}(y_c = 1|\mathbf{x})$ at each level of the hierarchy. Additionally, we report recall@ k at each level of the hierarchy. which indicates how many of all relevant classes were predicted in the set of k predictions coming from a classifier. Recall is formally defined as:

$$\text{recall@}k(\mathbf{y}, \mathbf{h}(x)) = \frac{1}{\|\mathbf{y}\|_1} \sum_{j \in \hat{\mathcal{Y}}_k} \mathbb{I}[y_j = 1] \quad (9)$$

where $\|\mathbf{y}\|_1$ is L_1 norm of vector \mathbf{y} – number of positive classes, $\hat{\mathcal{Y}}_k$ is a set of k labels with top conditional probabilities $\eta_c(\mathbf{x})$ predicted by \mathbf{h} for \mathbf{x} . We calculate recall@ k separately for each level. Notice that in this case, $\|\mathbf{y}\|_1$ is always equal to 1, as only one class is relevant per hierarchy level. Because of that, in this case, recall@1 is indifferent to accuracy. It is also easy to see that the optimal decision rule when predicting for recall@ k is to select top- k classes with the highest conditional probabilities $\eta_c(\mathbf{x})$ (Lapin, Hein, and Schiele 2018). This makes recall@ k align with our goal of measuring the quality of probability estimates while being more intuitive than logistic loss.

7 Results

7.1 Results for country specific classifier (trained on Polish language only)

Figure 5 compares the accuracy of different classification approaches for job codes across various levels of the KZIS occupation classification system, ranging from 1-digit to 6-digit codes, with ISCO classification at the 4-digit level. Across all datasets (Overall, CBOP, and Hand-coded job offers), transformer-based models consistently outperform linear models. Larger transformer models (HerBERT-large and XLM-RoBERTa-large) achieve the highest accuracy scores across all levels and datasets but are slightly better than base models.

Notably, the top-down approach consistently yields better results than the bottom-up ap-

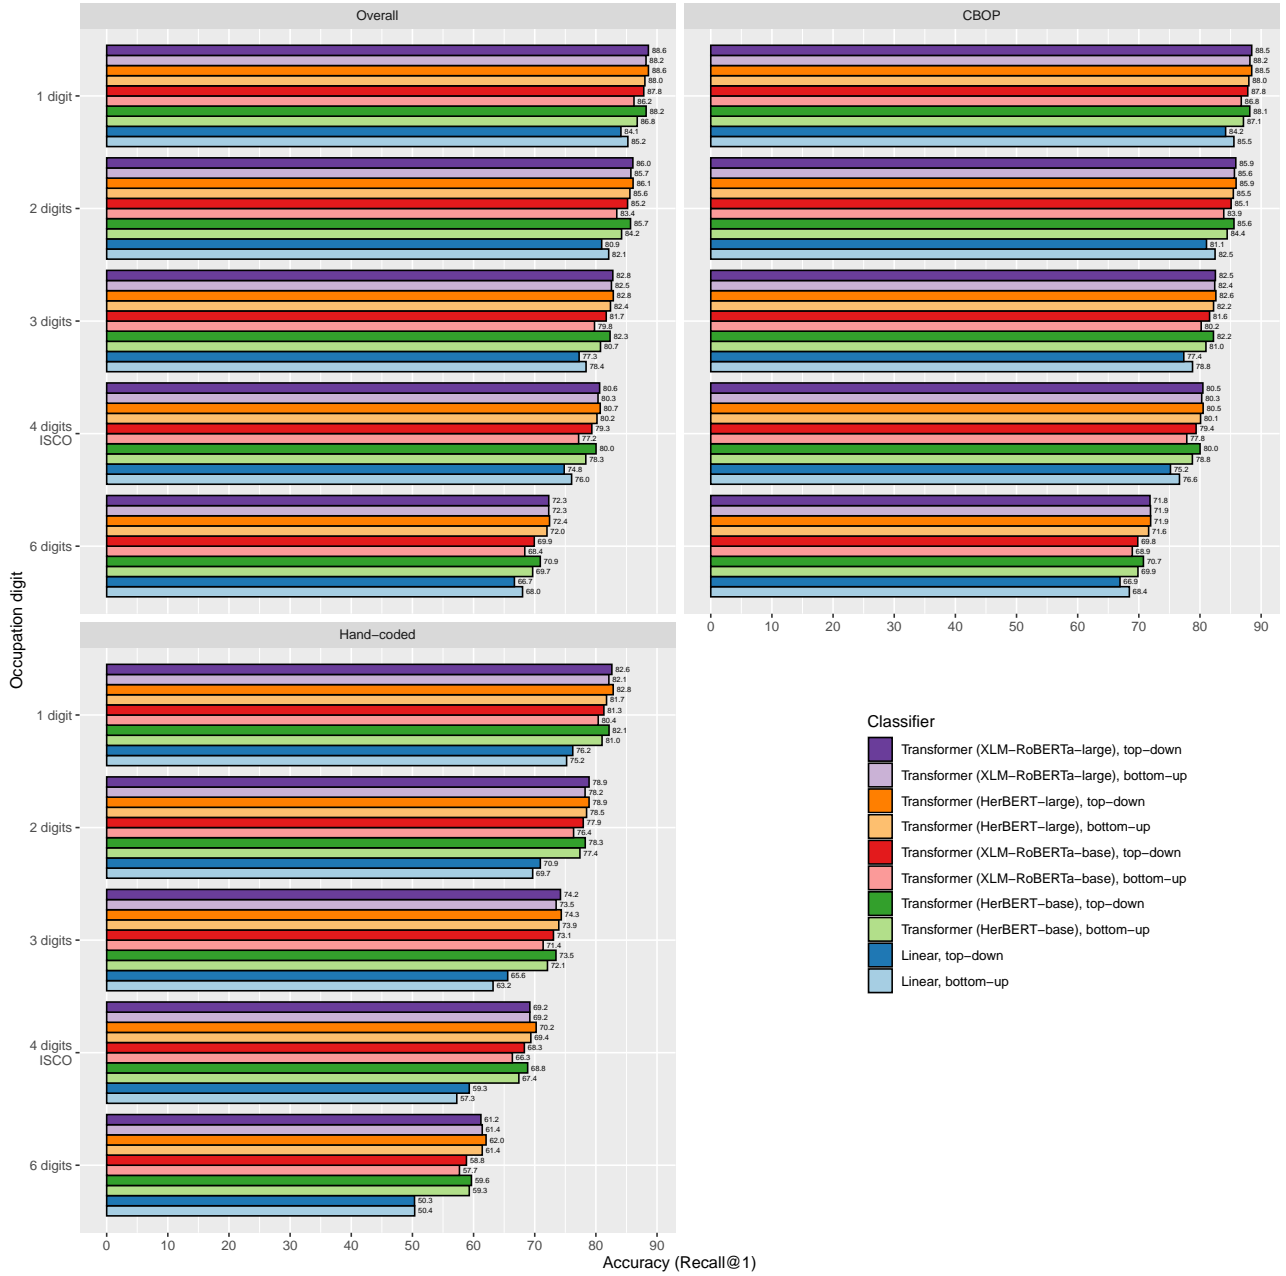


Figure 5: Results for the recall@1 classification of Polish job advertisements by classifier

proach for all models and coding levels. This improvement is particularly pronounced for hand-coded data, indicating that taking the hierarchical structure of occupational codes into account significantly enhances prediction accuracy. For instance, for the hand-coded dataset for HerBERT-base the difference at ISCO level is close to 1.5p.p. The same relation is present for the linear model where the differences between top-down and bottom-up are close to 2 p.p. (e.g., ISCO level)

As expected, the accuracy of all models decreases as the level of detail in the occupational codes increases. For example, the best-performing model (XLM-RoBERTa-large and HerBERT-large with top-down approach) achieves around 88.6% accuracy for 1-digit codes but drops to about 72.3% for 6-digit codes. This trend is consistent across all models and approaches, reflecting the increasing difficulty of classification as the categories become more specific.

Results for country-specific models suggest that large models perform similarly as base ones and country-specific BERT models performs slightly better than the multilingual BERT model. It should be noted that the XLM-RoBERTa-base model consist of 279M parameters while HerBERT 110M so comparison between these models it not entirely fair. However, it suggest that a smaller model for a specific language may perform similarly as larger model trained on multi-language dataset. In the next section, we will compare the multilingual model where the dataset contains two languages – Polish and English.

7.2 Results for bilingual classifier (trained on Polish and English languages)

7.2.1 Overall results

Figure 6 compares the accuracy (measured as Recall@1) of different classifiers for occupational coding across various levels of detail, from 1-digit to 6-digit codes, as well as ISCO codes. The classifiers include different transformer models (HerBERT and XLM-RoBERTa) using both top-down and bottom-up approaches. The plot also compares these automated methods to hand-coded results.

The results show that across all levels of occupational coding detail, the models using Google Translate consistently outperform those using Argos Translate. This difference is particularly noticeable for more detailed coding levels. For example, at the 6-digit level, the best-performing

Google model (Polish Transformer XLM-RoBERTa-base, bottom-up) achieves about 88% accuracy, while the corresponding Google Translate model achieves only about 86.4% accuracy. Notably, the HerBERT model, which was specifically designed for Polish, performs poorly when applied to English text, highlighting the importance of using language-appropriate models.

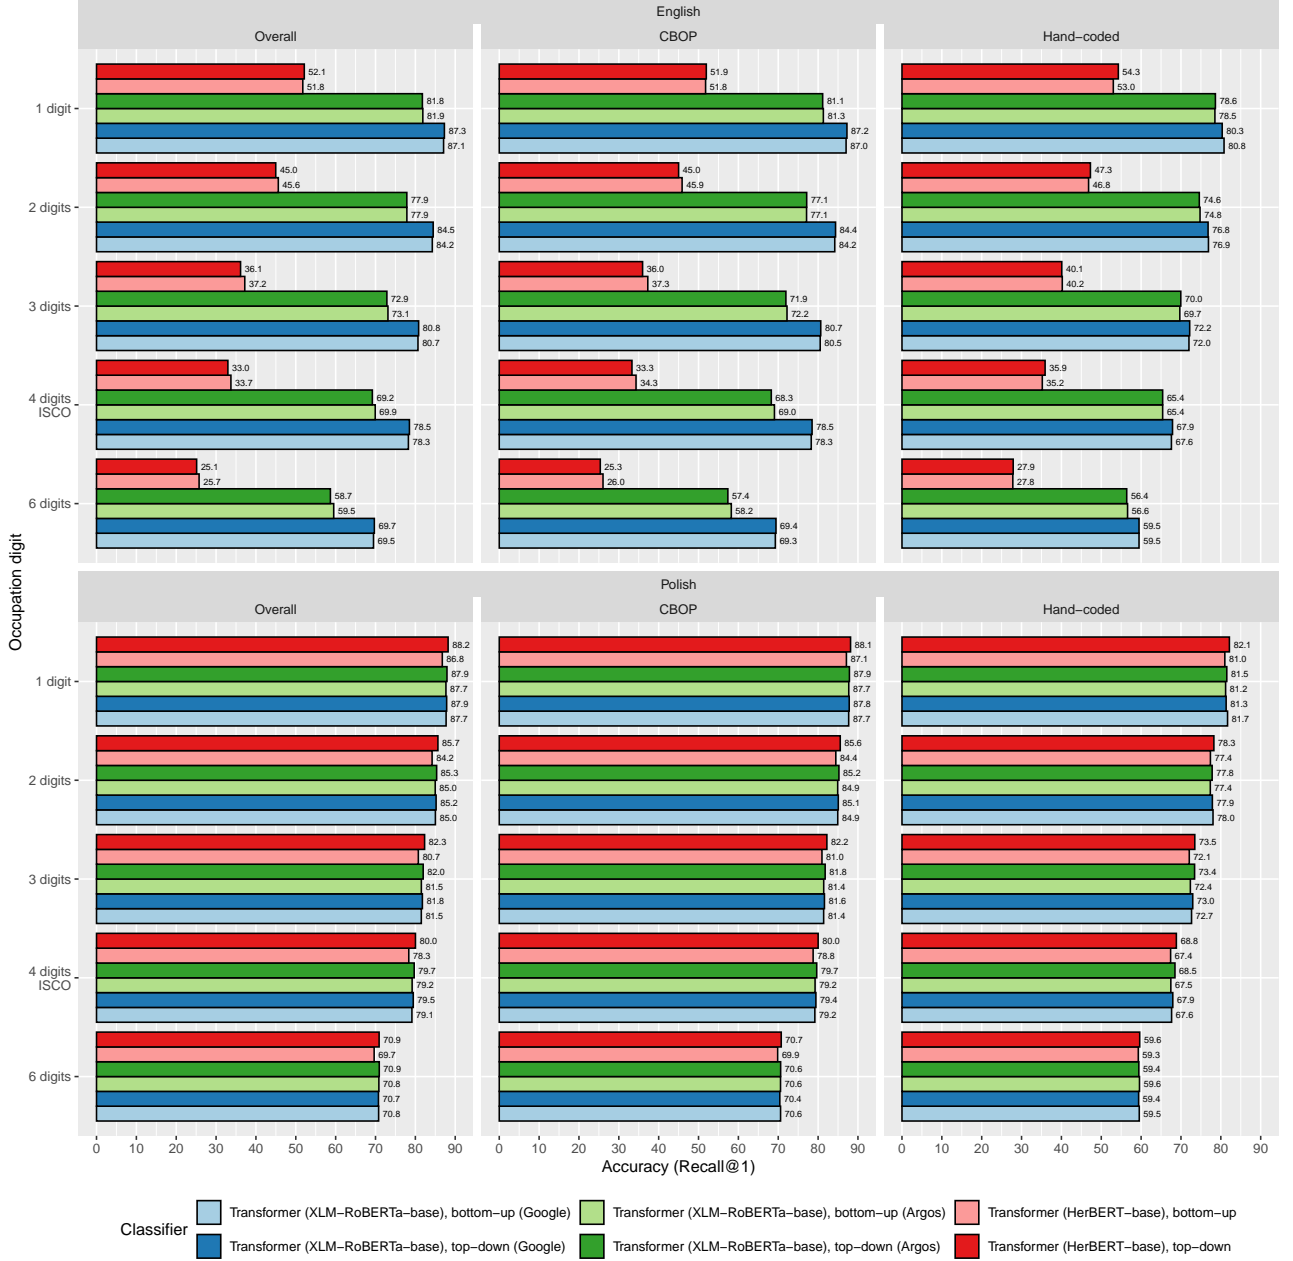


Figure 6: Results for the recall@1 classification of Polish job advertisements by multilingual classifier (XLM-RoBERTa-base)

Interestingly, the Polish versions of the models tend to perform slightly better than their English counterparts for both Argos and Google translations. This suggests that maintaining the original language (Polish in this case) or using a closely related language for translation might preserve more relevant information for occupational coding tasks. The poor perfor-

mance of HerBERT on English text further supports this observation, emphasizing the need for language-specific or truly multilingual models in cross-lingual tasks.

As observed in the previous plot, the top-down approach generally yields slightly better results than the bottom-up approach, although the difference is less pronounced here. The automated methods significantly outperform hand-coding across all levels of detail, with the gap widening for more granular classification tasks. The substantial difference in performance between Argos and Google translations highlights the importance of choosing appropriate translation tools in multilingual occupational coding tasks, with the closed-source Google Translate showing superior results in this context. This finding, combined with the language-specific performance differences, underscores the complexity of cross-lingual occupational coding and the need for careful selection of both translation tools and language models.

7.2.2 Detailed results for the hand-coded data

This section presents a more detailed analysis of the combined hand-coded data set, comprising 10,000 and 1,000 records. The data primarily come from online sources, and the classifier will be employed for further analysis of this type of data. Table 7.1 contains the corresponding data for the transformer model. The results are presented separately for the training and test data. For the training dataset, the transformer model achieves over 92% agreement for each group, while the logistic model performs significantly worse for the 1, 3, 4, and 9 groups.

In order to evaluate the efficacy of a given approach, it is essential to consider the results of the tests conducted. Once again, the transformer model demonstrated superior performance for all groups, with a precision rate exceeding 68% (with the exception of Occupation Group 6). In contrast, the logistic model exhibited lower precision rates for some groups, with a difference of approximately 10 percentage points. For instance, the logistic model for Group 3 achieved a precision rate of 57.5%, while the transformer model surpassed this rate, reaching over 68%. These results indicate that the transformer model is the optimal choice. Furthermore, we emphasize that the model’s performance can be validated using out-of-sample data, particularly new advertisements and clerical reviews with expert coding.

The confusion matrices for the XLM-RoBERTa-base transformer model reveal its performance in occupational coding tasks across Polish and English languages. Overall, the model

Table 7.1: Confusion matrix for 1 digit Occupation Groups (codes) for the transformer model with top-down approach on the test data

Country specific model – HerBERT										
Base model (110M parameters)										
True	1	2	3	4	5	6	7	8	9	N
1	81.0	11.6	2.6	1.7	2.2	0.0	0.9	0.0	0.0	232
2	3.1	89.1	4.4	1.5	1.6	0.0	0.3	0.0	0.0	1,273
3	2.6	15.7	67.2	2.6	4.8	0.0	5.2	1.0	1.0	503
4	0.8	12.8	4.5	70.9	3.0	0.0	0.0	1.5	6.4	265
5	1.8	4.0	4.5	1.5	85.7	0.0	0.8	0.0	1.8	399
6	14.3	14.3	0.0	0.0	0.0	57.1	0.0	0.0	14.3	7
7	0.8	1.4	2.9	0.0	0.2	0.0	89.2	2.7	2.9	518
8	0.0	0.4	0.8	1.6	0.8	0.0	17.8	76.4	2.3	258
9	0.0	0.9	2.7	4.5	3.2	1.4	7.7	0.9	78.6	220
Large model (340M parameters)										
1	82.8	11.2	1.7	1.7	2.2	0.0	0.4	0.0	0.0	232
2	3.5	88.2	4.8	1.6	1.6	0.0	0.3	0.0	0.0	1,273
3	2.0	13.9	68.8	2.6	4.4	0.0	6.0	1.2	1.2	503
4	1.1	10.2	6.0	71.3	3.0	0.0	0.0	1.1	7.2	265
5	1.5	4.3	6.3	1.0	85.0	0.0	0.8	0.0	1.3	399
6	0.0	14.3	0.0	0.0	0.0	71.4	0.0	0.0	14.3	7
7	0.6	1.4	1.7	0.2	0.4	0.0	90.7	2.7	2.3	518
8	0.0	0.0	1.9	1.2	0.4	0.0	18.2	77.9	0.4	258
9	0.5	0.0	2.7	5.0	2.7	2.3	9.1	0.9	76.8	220
Multilingual model – XLM-RoBERTa-base										
Polish language										
	1	2	3	4	5	6	7	8	9	
1	76.7	13.8	5.2	1.7	1.7	0.0	0.9	0.0	0.0	232
2	3.5	87.0	5.3	2.0	1.6	0.0	0.5	0.0	0.0	1,273
3	2.2	14.1	67.0	3.2	6.4	0.0	4.6	1.4	1.2	503
4	1.1	10.9	6.0	69.8	1.9	0.0	0.0	2.6	7.5	265
5	1.5	3.5	6.5	1.3	85.2	0.0	0.3	0.0	1.8	399
6	0.0	14.3	0.0	0.0	0.0	57.1	0.0	0.0	28.6	7
7	0.8	1.4	3.1	0.0	0.2	0.0	87.3	4.1	3.3	518
8	0.0	0.8	0.4	1.6	0.4	0.0	19.0	76.7	1.2	258
9	0.5	0.5	2.3	6.4	3.2	1.4	7.3	1.4	77.3	220
English language										
1	78.9	14.1	3.0	1.5	1.5	0.0	1.0	0.0	0.0	232
2	3.9	85.1	6.9	2.0	1.6	0.0	0.6	0.0	0.0	1,273
3	2.4	12.6	67.8	3.9	5.0	0.0	5.4	0.9	2.0	503
4	1.6	10.5	4.3	69.9	1.6	0.0	0.8	3.9	7.4	265
5	1.3	4.2	9.1	1.8	81.3	0.0	1.0	0.3	1.0	399
6	0.0	14.3	0.0	0.0	0.0	57.1	0.0	0.0	28.6	7
7	0.6	1.2	3.1	0.2	0.2	0.0	87.6	3.9	3.3	518
8	0.0	0.4	0.8	1.6	0.4	0.0	17.3	77.5	2.0	258
9	0.5	0.5	1.8	6.4	1.8	0.9	9.2	1.4	77.5	220

demonstrates high accuracy, with the majority of predictions correctly falling on the diagonal for both languages. Notably, Occupation Groups 2, 5, and 7 consistently show over 80% correct classifications, indicating the model’s strength in identifying these occupational groups.

However, some common misclassification patterns emerge, such as Group 2 being mistaken for Group 1, and Group 3 for Group 2, suggesting potential similarities or overlaps in these occupational definitions that challenge the model’s discriminative capabilities.

Interestingly, while the overall performance is comparable between Polish and English, some nuanced differences are observed. For instance, Occupation Group 1 is more accurately classified in English (78.9%) compared to Polish (76.7%), whereas Group 5 shows superior performance in Polish (85.2%) over English (81.3%). These variations highlight the potential impact of language-specific nuances on occupational coding accuracy, emphasizing the importance of considering linguistic context in model development and application.

The matrices also reveal challenges in classifying low-frequency categories, exemplified by Occupation Group 6, which has only 7 samples and exhibits poor performance with a high misclassification rate to Group 9. This underscores the persistent challenge in machine learning of accurately handling rare or underrepresented classes, particularly in specialized domains like occupational coding.

Despite these challenges, the consistency in misclassification patterns across both languages suggests that many errors stem from inherent similarities between certain occupational categories rather than language-specific issues. This insight is valuable for future refinements of occupational classification systems and the development of more robust coding models. The strong overall performance of the XLM-RoBERTa-base model across languages demonstrates its potential as a powerful tool for automated occupational coding, while also highlighting areas for future improvement, particularly in distinguishing between similar categories and handling low-frequency occupations. Detailed results are presented in Supplementary Materials.

7.3 Results for multilingual dataset (24 languages)

In this section, we present the results for a multilingual dataset that has been trained with XLM-RoBERTa models. Table 7.2 presents the results for the test dataset only. The table also includes two approaches to modelling the hierarchy for 1, 2, 4 (ISCO) and 6 (KZiS) digits. The results are comparable to those observed for the Polish and bilingual models. The overall classification rate for 1-digit codes is approximately 84% (approximately 2-4 p.p. lower than the first two approaches), and the hand-coded classification rate is approximately 78%. As

previously noted, the inclusion of the hierarchy improves predictions for the hand-coded data as the hierarchy is traversed.

Table 7.2: Recall@1 metric for the multilingual test data by dataset and transformer approach

Dataset	Hierarchy	1 digits	2 digits	ISCO (4 digits)	KZiS (6 digits)
Overall	bottom-up	84.37	80.87	73.67	64.42
	top-down	84.40	80.70	73.25	63.74
CBOP	bottom-up	84.38	80.88	73.77	64.41
	top-down	84.43	80.75	73.46	63.90
Hand-coded	bottom-up	77.82	72.95	62.37	53.51
	top-down	78.59	73.62	62.46	53.15

As indicated in Section 5.2, the sample size for languages differs considerably in the test data. Accordingly, in order to evaluate the accuracy of the multilingual classifier, we present the results for each translated test data set, including the complete data set and the sample utilized for training. Figure 7 depicts the accuracy (recall@1) for each language across 1-digit, 4-digit (ISCO) and 6-digit (KZiS) occupation codes, employing both the bottom-up and top-down approaches. The figures illustrate the results for the complete and sampled hand-coded test data sets.

As expected, the highest classification was observed for English (translated with Google Translate, the most prevalent language in the XLM-RoBERTa dataset) and Polish (original), with a near 60% accuracy at 6-digit codes and a near 80% accuracy at 1-digit codes. The top-down method yielded improvements in classification for 67% of languages at the 1-digit level and 20% at the ISCO level. The top-down approach yielded an average improvement of 0.3 p.p. at the 1-digit code level, while exhibiting a decline of -0.3 p.p. at the 4-digit code level. Notably, this approach consistently enhanced classification outcomes for Polish and English languages, suggesting that the impact on other languages may be diminished due to potential discrepancies in translation from English to other languages via the Argos library. Moreover, the degree of accuracy at the 6-digit level exhibits a notable disparity, with a near 60% success rate for English and a figure approaching 50% for Irish. This trend persists at the ISCO level, while the variance at the 1-digit level is less pronounced.

The results for the sampled part of the test data diverge significantly from the results for the entire dataset. For instance, the French result for 6-digit codes was approximately 40%, while the overall test data yielded a figure exceeding 55%. Conversely, the Polish and English results




Figure 7: Results for the multilingual model for 1-digit, 4-digits and 6-digit codes by the approach (bottom-up and top-down) for hand-coded test data (whole and sampled). Note X axis is sorted according to the results for the whole data.

exhibited a reversed pattern, with values of over 80% and 90%, respectively, in comparison to approximately 80% for both.

A modest correlation was observed between the accuracy of the XLM-Roberta model and the coverage of languages. This is evidenced by the Pearson’s and Spearman’s correlation coefficients between accuracy and the share of tokens, which were 0.26 and 0.17, respectively (see Table 5.8). The highest Pearson’s and Spearman’s correlation coefficients were observed for the share of cases in the test data (0.37 and 0.22, respectively). The detailed results for this classifier are presented in the Supplementary Materials.

8 Discussion

In this paper, we propose a classifier and conditional probability estimator of occupation codes for job vacancy statistics that takes into account the hierarchical structure of the occupation classification. Occupations as proxies of job-related skills seem to be especially important in

explaining the technological, demographic, and social transformation in the labour markets. The proposed classifier can be used for automated occupational coding by official statistics or labour market institutions that collect unstructured information about labour demand and formulate policies. To our knowledge this is the first publicly available model for 24 languages along with detailed description on the data and predictions quality.

Our approach can be used with information-abundant online and administrative data to enrich survey-based job vacancy data. This is especially important nowadays, when survey research faces the problems of non-response and under-reporting. The use of online job advertisements is the main job search method and often also recruitment. They contain plenty of unstructured information. While statistical offices lack methods to classify such text, we provide one. Our approach is dedicated to job offers, and our procedure shows how to incorporate country-specificity on the example of Poland.

We use the Polish Classification of occupations and specializations, which is six digits, going one level further than the international ISCO classification, and is similar to the ESCO classification. We compare models trained on English and Polish languages and show the limitations of applying generic methods based on English language. We find that the country-specific model performs better than the multilingual one. This confirms the need for language-specific models, as it outperformed the generic ones that were only prepared for our task, and not trained with Polish language. However, there were examples of occupations (e.g., managers) that were well classified in English. This might result from the fact that some occupational titles have English origin and are often used in English form in job advertisements. This may be an example for other language, and institution-specific countries.

In this paper, our focus is on the Polish KZiS and ISCO. However, it should be noted that other classifications may be of interest to users, including O*NET, ESCO and SOC. One potential avenue for investigation is the utilisation of publications that align with one another in terms of their respective classifications. For example, Eurostat has published a report entitled "The crosswalk between ESCO and O*NET" (see <https://esco.ec.europa.eu/en/about-esco/data-science-and-esco/crosswalk-between-esco-and-onet>), which may be used to align our predictions for ISCO with the O*NET classification system. An alternative approach may be to draw upon the work of Stęchły and Woźniak (2023), who establishes a link between

the Polish KZiS and ESCO classifications. Ultimately, researchers may utilise the pre-trained models presented here and develop their own classifiers.

Acknowledgements

We would like to thank Natalie Shlomo and Piotr Chlebicki for their valuable feedback on the previous draft of this paper.

Cherniaiev Herman, Pater Robert, Wydmuch Marek work and research were conducted in the Educational Research Institute (IBE) in Warsaw, and funded by the Polish Ministry of Education and Science within a project *Supporting IQS functioning and improvements in order to use its solutions in achieving the country's development strategy aims*, ZSK5 (POWR.02.13.00-IP.02-00-002/20). The preparation of hand-coded datasets was financed by IBE. Maciej Beręsewicz's work was partially supported by the Polish National Agency for Academic Exchange (NAWA) under the project number BPN/BEK/2023/1/00099.

Marek Wydmuch and Maciej Beręsewicz contributed equally to the paper. Marek Wydmuch was responsible for developing and coding the classifier, Maciej Beręsewicz was responsible for preparing training and testing datasets, computation, preparation of the manuscript, and assessing the results. Herman Cherniaiev collected and initially processed the data on job offers. Robert Pater reviewed the literature and provided the economic and the educational context of research.

Information on software and models

- Software: <https://github.com/OJALAB/job-ads-classifier>
- Models: <https://repod.icm.edu.pl/dataset.xhtml?persistentId=doi:10.18150/OCUTSI>
- Tutorial: <https://colab.research.google.com/drive/1a425aagT0lczRxXPWoUlf5aFxFxUII37nh?usp=sharing>

Notes

¹See <https://www.cedefop.europa.eu/en/tools/skills-online-vacancies/occupations/>

²<https://www.gov.pl/web/edukacja/zawody-szkolnictwa-branzowego> (as of 2024.07.31).

³Available at <https://psz.praca.gov.pl/rynek-pracy/bazy-danych/klasyfikacja-zawodow-i-specjalnosci/wyszukiwarka-opisow-zawodow> (as of 2024.07.18).

⁴<https://psz.praca.gov.pl/rynek-pracy/bazy-danych/infodoradca> (as of 2024.07.18).

⁵Available at <https://stat.gov.pl/Klasyfikacje/doc/kzs/slownik.html> (as of 2024.07.18)

⁶The value of this crosswalk is providing occupation descriptions from the ESCO classification and connecting occupations with skills, what has not been done in the KZiS. To connect KZiS with ESCO authors used an expert method supported with language modelling and semantic similarity based on Sentence-BERT model. They determined the relations for all KZiS occupations by matching existing ESCO occupations to KZiS. In linking both classifications, the authors used several types of semantic relations: exact match, broader match, narrower match, and close match. They excluded false friend (no) matches and weak matches. Available at https://esco.ec.europa.eu/en/classification/occupation_main (as of 2024.07.18)

⁷The following codes were identified 121101, 121904, 214921, 242208, 242211, 242213, 261103, 315202, 315209, 325504, 334102, 421402.

⁸For an illustrative example, please refer to this website: <https://warszawa.praca.gov.pl/zgloszenie-oferty-pracy> (in Polish). In Supplementary Materials we provide an example paper form in Polish and English translated using DeepL software.

⁹See <https://www.gov.pl/web/edukacja/prognoza-zapotrzebowania-na-pracownikow-w-zawodach-szkolnictwa-branzowego-na-krajowym-i-wojewodzkiem-rynku-pracy-2024>

¹⁰More information is available at the following website <https://oferty.praca.gov.pl/portal/index.cbop#/dlaInt>

¹¹The codes used for this purpose are available for download online at <https://github.com/OJALAB/CBOP-datasets>

¹²The five most prevalent codes were: The following job titles were identified: 522301 "Sprzedawca" (Sales Assistant), 432103 "Magazynier" (Warehouse Operator), 931301 "Pomocniczy robotnik budowlany" (Construction Worker's Helper), 911207 "Pracownik utrzymania czystości" (Cleaner), 515303 "Robotnik gospodarczy" (Maintenance Worker).

¹³Information about the number is presented in Table A.1 in Supplementary Materials.

¹⁴The number of words was counted using `stri_count_words` function from the `stringi` package and number of characters were counted using `nchar` from the base R

¹⁵it should be noted that for 4 ads we also had information from CBOP

¹⁶Note that this number is larger than the number of codes because for certain codes more than one case is provided.

¹⁷Detailed information: 1 group: 9, 2 group: 84, 3 group: 46, 4 group: 1, 5 group: 5, 6 group: 1, 7 group: 27,

8 group: 8, 9 group: 5). List of these codes is provided here <https://github.com/OJALAB/job-ads-datasets/blob/main/data/codes-not-coveted.csv>.

References

- Babbar, R., and B. Schölkopf. 2017. “Dismec: Distributed sparse machines for extreme multi-label classification.” In *Proceedings of the tenth ACM international conference on web search and data mining*, 721–729. New York, NY, USA: Association for Computing Machinery.
- Barbedo, J. G. s., and A. Lopes. 2006. “Automatic genre classification of musical signals.” *EURASIP Journal on Advances in Signal Processing* 2007:1–12.
- Beręsewicz, M., H. Cherniaiev, and R. Pater. 2021. “Estimating the number of entities with vacancies using administrative and online data.” *arXiv preprint arXiv:2106.03263*.
- Beręsewicz, M., and R. Pater. 2021. *Inferring job vacancies from online job advertisements*. Publications Office of the European Union.
- Beręsewicz, M. E., G. Białkowska, K. Marcinkowski, M. Maślak, P. Opiela, P. Katarzyna, and R. Pater. 2021. “Enhancing the Demand for Labour survey by including skills from online job advertisements using model-assisted calibration.” *Survey Research Methods* 15 (2): 147–167. <https://doi.org/10.18148/srm/2021.v15i2.7670>. <https://ojs.ub.uni-konstanz.de/srm/article/view/7670>.
- Beygelzimer, A., J. Langford, Y. Lifshits, G. Sorkin, and A. Strehl. 2009. “Conditional Probability Tree Estimation Analysis and Algorithms.” In *Proceedings of the Twenty-Fifth Conference on Uncertainty in Artificial Intelligence*, 51–58. Arlington, Virginia, USA: AUAI Press.
- Bhatia, K., H. Jain, P. Kar, M. Varma, and P. Jain. 2015. “Sparse Local Embeddings for Extreme Multi-label Classification.” In *Advances in Neural Information Processing Systems 28*, 730–738. Curran Associates, Inc.
- Boselli, R., M. Cesarini, F. Mercorio, and M. Mezzanzanica. 2017. “Using machine learning for labour market intelligence.” In *Machine Learning and Knowledge Discovery in Databases: European Conference, ECML PKDD 2017, Skopje, Macedonia, September 18–22, 2017, Proceedings, Part III 10*, 330–342. Springer.

- Chen, H., and C. M. Mason. 2024. “Explainable AI (XAI) for Constructing a Lexicon for Classifying Green Energy Jobs: A Comparative Analysis of Occupation, Industry and Location Composition with Traditional Energy Jobs.” *IEEE Access*.
- Colombo, E., F. Mercorio, and M. Mezzanzanica. 2018. “Applying machine learning tools on web vacancies for labour market and skill analysis.” *Terminator or the Jetsons*.
- Conneau, A., K. Khandelwal, N. Goyal, V. Chaudhary, G. Wenzek, F. Guzmán, É. Grave, M. Ott, L. Zettlemoyer, and V. Stoyanov. 2020. “Unsupervised Cross-lingual Representation Learning at Scale.” In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, 8440–8451.
- Dembczynski, K., W. Kotłowski, W. Waegeman, R. Busa-Fekete, and E. Hüllermeier. 2016. “Consistency of probabilistic classifier trees.” In *ECML PKDD 2016 : machine learning and knowledge discovery in databases*, 9852:511–526. Lecture Notes in Computer Science. Springer.
- Devlin, J. 2018. “Bert: Pre-training of deep bidirectional transformers for language understanding.” *arXiv preprint arXiv:1810.04805*.
- European Commission. 2022. *Machine Learning Assisted Mapping of Multilingual Occupational Data to ESCO*. Technical report. European Commission.
- Falcon, W., and T. P. L. Team. 2019. *PyTorch Lightning*. <https://www.pytorchlightning.ai>.
- Fox, J. 1997. *Applied regression analysis, linear models, and related methods*. Sage.
- Gałęcka-Burdziak, E., and R. Pater. 2015. “Ile jest wolnych miejsc pracy w Polsce?” *Gospodarka Narodowa. The Polish Journal of Economics* 279 (5): 171–186.
- Hersh, J., and M. Harding. 2018. “Big data in economics.” *IZA World of Labor*, no. 451.
- Hershbein, B., and L. B. Kahn. 2018. “Do recessions accelerate routine-biased technological change? Evidence from vacancy postings.” *American Economic Review* 108 (7): 1737–72.
- Jasinska-Kobus, K., M. Wydmuch, K. Dembczynski, M. Kuznetsov, and R. Busa-Fekete. 2020. *Probabilistic Label Trees for Extreme Multi-label Classification*. arXiv: 2009.11218 [cs.LG].

- Klein, G., Y. Kim, Y. Deng, J. Senellart, and A. Rush. 2017. “OpenNMT: Open-Source Toolkit for Neural Machine Translation.” In *Proceedings of ACL 2017, System Demonstrations*, 67–72. Vancouver, Canada: Association for Computational Linguistics, July. <https://www.aclweb.org/anthology/P17-4012>.
- Lapin, M., M. Hein, and B. Schiele. 2018. “Analysis and Optimization of Loss Functions for Multiclass, Top-k, and Multilabel Classification.” *IEEE Transactions on Pattern Analysis and Machine Intelligence* 40 (7): 1533–1554. <https://doi.org/10.1109/TPAMI.2017.2751607>.
- Loshchilov, I., and F. Hutter. 2017. “Decoupled weight decay regularization.” *arXiv preprint arXiv:1711.05101*.
- Lumley, T. 2004. “Analysis of Complex Survey Samples.” R package version 2.2, *Journal of Statistical Software* 9 (1): 1–19.
- Ministry of Family, Labour and Social Policy. 2014. *Rozporządzenie Ministra Pracy i Polityki Społecznej z dnia 7 sierpnia 2014 r. w sprawie klasyfikacji zawodów i specjalności na potrzeby rynku pracy oraz zakresu jej stosowania*. Tekst jednolity Dz. U. z 2018 r. poz. 227 z późn. zm. (Dz.U. z 2021 r. poz. 2285 oraz Dz.U. z 2022 r. poz. 853).
- Morin, F., and Y. Bengio. 2005. “Hierarchical Probabilistic Neural Network Language Model.” In *Proceedings of the Tenth International Workshop on Artificial Intelligence and Statistics*, 246–252. Society for Artificial Intelligence / Statistics.
- Mroczkowski, R., P. Rybak, A. Wróblewska, and I. Gawlik. 2021. “HerBERT: Efficiently Pre-trained Transformer-based Language Model for Polish.” In *Proceedings of the 8th Workshop on Balto-Slavic Natural Language Processing*, 1–10. Kiyv, Ukraine: Association for Computational Linguistics, April. <https://www.aclweb.org/anthology/2021.bsnlp-1.1>.
- Ooms, J. 2024. *cld3: Google’s Compact Language Detector 3*. [10.32614/CRAN.package.cld3](https://cran.r-project.org/web/packages/cld3/index.html). <https://github.com/ropensci/cld3>.

- Paszke, A., S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, et al. 2019. “PyTorch: An Imperative Style, High-Performance Deep Learning Library.” In *Advances in Neural Information Processing Systems 32*, edited by H. Wallach, H. Larochelle, A. Beygelzimer, F. d’Alché-Buc, E. Fox, and R. Garnett, 8024–8035. Curran Associates, Inc. <http://papers.neurips.cc/paper/9015-pytorch-an-imperative-style-high-performance-deep-learning-library.pdf>.
- Radzikowski, M. 2023. “Monitorowanie działalności gospodarczej. Gospodarka po pandemii COVID-19.” Chap. 11. Oficyna Wydawnicza SGH.
- Silla, C. N., and A. A. Freitas. 2011. “A survey of hierarchical classification across different application domains.” *Data mining and knowledge discovery* 22:31–72.
- Stęchły, W., and I. Woźniak. 2023. “PROPOZYCJA POWIĄZANIA ZAWODÓW Z POLSKIEJ KLASYFIKACJI ZAWODÓW I SPECJALNOŚCI Z ZAWODAMI ESCO.” *Rynek Pracy* 187 (4): 22–45.
- Sun, A., and E.-P. Lim. 2001. “Hierarchical text classification and evaluation.” In *Proceedings 2001 IEEE International Conference on Data Mining*, 521–528. IEEE.
- Tillé, Y., and A. Matei. 2023. *sampling: Survey Sampling*. [10.32614/CRAN.package.sampling](https://CRAN.R-project.org/package=sampling). <https://CRAN.R-project.org/package=sampling>.
- Turrell, A., B. Speigner, D. Copple, J. Djumalieva, and J. Thurgood. 2021. “Is the UK’s productivity puzzle mostly driven by occupational mismatch? An analysis using big data on job vacancies.” *Labour Economics* 71:102013.
- Turrell, A., B. Speigner, J. Djumalieva, D. Copple, and J. Thurgood. 2022. “6. Transforming Naturally Occurring Text Data into Economic Statistics.” In *Big Data for Twenty-First-Century Economic Statistics*, 173–208. University of Chicago Press.
- Vaswani, A. 2017. “Attention is all you need.” *arXiv preprint arXiv:1706.03762*.
- Wolf, T. 2019. “Huggingface’s transformers: State-of-the-art natural language processing.” *arXiv preprint arXiv:1910.03771*.

Supplementary materials for the paper

Multilingual hierarchical classification of job advertisements for job vacancy statistics

List of Supplementary Materials

1. Additional information on the data and its quality (in PDF)
2. Models results in Excel Spreadsheet: `models-accuracy.xlsx`.
3. A short tutorial on using the software and models for classification and estimation of conditional probability.

A Additional information on the data and its quality

Table A.1: Number of ads web-scraped between 2021-06-28 and 2021-11-08

	Source	N
1	www.pracuj.pl	447,624
2	www.praca.egospodarka.pl	330,465
3	www.praca.pl	303,771
4	www.aplikuj.pl	228,575
5	www.nuzle.pl	111,623
6	jobdesk.pl	106,119
7	gratka.pl	91,935
8	CBOP	61,596
9	www.infopraca.pl	45,164
10	www.gowork.pl	24,962
11	www.jobs.pl	22,741
12	www.karierawfinansach.pl	12,774
13	www.jober.pl	9,139
14	www.absolvent.pl	2,744
15	nofluffjobs.com	2,172
16	www.goldenline.pl	2,130
17	praca.dlastudenta.pl	2,019
18	students.pl	414

Table A.2: Information on stratification variables and population size (in thousands)

	[1,10]	(10,25]	(25,50]	(50,100]	(100,150]	(150,200]	(200,300]	(300,Inf]	Σ
biurowy (<i>office</i>)	1.15	2.04	1.98	6.28	3.46	1.50	1.83	0.90	19.15
doradcy (<i>advisors</i>)	0.24	0.34	0.50	4.75	7.46	10.25	9.68	2.07	35.30
magazynier (<i>warehouseman</i>)	0.94	1.88	2.43	8.59	7.69	4.83	4.99	1.90	33.26
kucharze (<i>cooks</i>)	1.27	2.73	2.83	6.97	3.10	0.75	0.75	0.23	18.63
produkcja (<i>production</i>)	0.60	1.49	1.55	5.10	5.29	4.21	4.42	2.02	24.67
przedstawiciele (<i>representatives</i>)	0.26	0.43	0.51	4.97	10.74	7.33	6.76	1.84	32.84
specjaliści (<i>specialists</i>)	1.27	2.15	2.62	11.53	25.73	27.40	32.59	26.82	130.11
sprzątacze (<i>cleaners</i>)	0.56	1.40	1.24	2.60	0.70	0.54	0.00	0.00	7.04
sprzedawcy (<i>sellers</i>)	1.97	4.20	5.59	16.78	11.03	5.52	7.79	1.85	54.73
inne (<i>other</i>)	31.74	53.37	60.60	186.46	200.82	169.66	208.79	123.80	1,035.24
Σ	40.01	70.04	79.86	254.03	276.01	232.00	277.60	161.43	1,390.98

Table A.3: Sample size by source

Source	N	%
pracuj.pl	2,201	22.0
praca.egospodarka.pl	1,967	19.7
praca.pl	1,881	18.8
aplikuj.pl	880	8.8
nuzle.pl	764	7.6
jobdesk.pl	757	7.6
gratka.pl	561	5.6
CBOP	241	2.4
inne	750	7.5
Total	10,002	100

Table A.4: Phrases used, sample and population number of ads

Phrase	Sample (n)	Population (N)
blockchain block chain (<i>blockchain/block chain</i>)	10	39
cloud cloud computing chmura (<i>cloud/cloud computing</i>)	100	1,007
cyberbezpiecz (<i>cybersecurity</i>)	100	488
diagnosta laboratoryjny (<i>laboratory diagnostician</i>)	10	46
etyki biznesu (<i>business ethics</i>)	9	9
farmaceuta (<i>pharmacist</i>)	10	168
(frontend front end) developer (<i>(frontend/front end) developer</i>)	100	287
informatyk(a i) przemysł(a ej) industrial computer science (<i>industrial computer science</i>)	30	30
internet rzeczy IoT internet of things (<i>IoT/internet of things</i>)	14	14
kryminaln cyberprzest (<i>criminal/cybercriminals</i>)	22	22
lekarz weterynarii (<i>veterinarian</i>)	10	33
pianino fortepian pianina (<i>piano</i>)	4	4
położna (<i>midwife</i>)	10	30
robotyk (<i>roboticist/robotics specialist</i>)	100	1667
spawacz (mag mit mig) (<i>welder (mag/mit/mig)</i>)	10	219
tłumacz przysięgły (<i>sworn translator</i>)	5	5
toksykolog (<i>toxicologist</i>)	14	14
uczenie maszynowe machine learning data scientist (<i>machine learning/data scientist</i>)	150	354
virtual reality wirtualna rzeczywistość (<i>virtual reality</i>)	5	7
unspecified phase (the rest)	287	227,122
–	1,000	227,825

Table A.5: Unique 6- and 1-digit codes by strata for hand-coded 1k dataset

Stratum (Phrase)	Unique codes	
	6-digits	1-digit
other	138	8
robotyk (<i>roboticist/robotics specialist</i>)	30	6
cyberbezpiecz (<i>cybersecurity</i>)	29	3
cloud cloud computing chmura (<i>cloud/cloud computing</i>)	24	4
uczenie maszynowe machine learning data scientist (<i>machine learning/data scientist</i>)	21	5
informatyk(a i) przemysłow(a ej) industrial computer science (<i>industrial computer science</i>)	11	4
blockchain block chain (<i>blockchain/block chain</i>)	9	3
toksykolog (<i>toxicologist</i>)	8	3
etyki biznesu (<i>business ethics</i>)	8	4
lekarz weterynarii (<i>veterinarian</i>)	6	2
internet rzeczy IoT internet of things (<i>IoT/internet of things</i>)	6	2
kryminaln cyberprzest (<i>criminal/cybercriminals</i>)	6	1
diagnosta laboratoryjny (<i>laboratory diagnostician</i>)	5	2
farmaceuta (<i>pharmacist</i>)	5	3
(frontend front end) developer (<i>(frontend/front end) developer</i>)	4	1
położna (<i>midwife</i>)	4	3
pianino fortepian pianina (<i>piano</i>)	4	2
virtual reality wirtualna rzeczywistość (<i>virtual reality</i>)	3	2
spawacz (mag mit mig) (<i>welder (mag/mit/mig)</i>)	3	1
tłumacz przysięgły (<i>sworn translator</i>)	1	1

Note: Table A.5 provides information on the number of unique 6-digit and 1-digit codes for a given stratum defined by phrase. This can provide information about the quality of the initial stratification based on keyword searches. For example, the main codes for 'pianino|fortepian|pianina' were related to kitchen and restaurant organisation. The reason for this is that these keywords are related to the name of the restaurant rather than the piano itself.

Table A.6: Number of words by source and main groups

Source	Group	Min	Q1	Med	Mean	Q3	Max	N
CBOP	0	11	26	44	55.2	52	208	9
	1	2	38	66	99.8	118	2,561	3,043
	2	2	27	46	74.9	84	2,310	22,920
	3	1	29	50	77.0	89	1,406	16,152
	4	1	24	42	73.0	74	3,385	17,144
	5	1	19	31	40.3	49	3,215	26,434
	6	2	18	29.5	35.8	47	283	1,184
	7	1	19	30	40.1	50	1,174	34,769
	8	2	22	34	42.3	54	1,241	18,243
	9	1	19	31	40.0	50	1,067	27,346
ESCO	0	4	6.2	18	44.5	56.2	138	4
	1	1	4	22	67.9	88.5	1,233	156
	2	1	4	30	63.8	88	1,035	681
	3	1	4	32	61.2	99	634	425
	4	1	3	24	37.7	67.5	222	67
	5	1	3	14	48.3	77	321	176
	6	1	6	18	51.8	52	283	31
	7	1	3	27	48.5	71.5	396	290
	8	1	5	21	47.8	65	209	210
	9	2	4	16	41.3	54.2	282	74
Thesaurus	0	10	22	34	28.7	38	42	3
	1	2	7	9	11.9	14	153	173
	2	2	5	6	8.9	10	54	209
	3	2	6	10	13.5	16	83	293
	4	2	4	6	8.5	8.5	38	19
	5	2	4	7	9.2	12.2	27	48
	6	3	5.8	10.5	12.1	17.2	32	36
	7	2	7	11	15.4	18	162	222
	8	4	11	19	27.1	34	232	319
	9	2	4	5	7.7	7.2	39	20
Hand	1	12	156	202	230.9	263	1,279	577
	2	11	151	217	255.5	289	2,836	3,113
	3	4	128	185	221.0	259	2,792	1,346
	4	6	101	149	191.6	227	2,795	821
	5	8	100	144	165.9	218	2,848	1,244
	6	15	68	90.5	119.7	163	338	18
	7	6	85.2	131	163.5	205	1,205	1,486
	8	6	76	134	156.6	197	1,459	713
	9	8	54	111	134.7	167.5	1,686	674
Hand 1k	1	78	108.5	145	153.9	165	495	59
	2	31	105.8	137	152.8	183	657	748
	3	48	92.5	113	121.0	137	326	108
	4	22	52.8	75	96.4	106.8	400	22
	5	39	61.2	78	84.8	95.2	192	30
	7	26	51	77	85.0	109	187	47
	8	23	53	66	74.9	95.5	160	15

(Continued on next page)

(Continued from previous page)

Source	Group	Min	Q1	Med	Mean	Q3	Max	N
	9	53	62.2	69.5	76.5	95.5	103	6
INFO+	1	16	97	171	237.0	308.5	1,260	91
	2	17	95	183.5	250.3	329	2,107	1,344
	3	13	85	173	227.8	307	1,374	1,533
	4	14	73.8	141	197.0	273	1,245	336
	5	15	82	151	210.5	288	1,355	714
	6	14	83	181	249.3	358.8	1,229	308
	7	8	78	144.5	225.4	289	1,924	1,540
	8	16	82	138.5	227.7	285	1,318	1,106
	9	22	75.5	110	176.3	209.5	988	35
KPRM	1	84	232.2	304.5	316.0	380.5	771	562
	2	62	147	189	208.1	252	852	1,924
	3	55	181	221	233.2	274	703	411
	4	82	126.2	170.5	178.6	230	300	44
Official	0	4	18.2	40	59.8	77.2	170	12
	1	1	6	28	70.0	110	410	708
	2	1	5	37	67.8	120.8	508	2,502
	3	1	4	29	59.9	95	574	1,961
	4	1	3	29	60.4	117	287	257
	5	1	4	27	58.3	101	307	521
	6	1	4	21	55.4	107.5	275	192
	7	1	4	32	60.5	114	327	1,466
	8	1	6	32	59.5	104	340	1,155
	9	1	4	17	39.3	58.8	396	426

Min = Minimum, Q1 = 1st Quartile, Med = Median,
Q3 = 3rd Quartile, Max = Maximum, N = Number of observations