

models_tutorial_classifier_2025

March 19, 2025

0.1 README

This is an example tutorial on how to use pre-trained models to predict occupation codes for online advertisements.

Note that the software allows you to use these models to fine tune them for your training and test data.

Models: <https://repod.icm.edu.pl/dataset.xhtml?persistentId=doi:10.18150/OCUTSI>

Publication: Beręsewicz, M., Wydmuch, M., Cherniaiev, H., & Pater, R. (2024). Multilingual hierarchical classification of job advertisements for job vacancy statistics. arXiv preprint arXiv:2411.03779.

How to cite:

```
@misc{beręsewicz2024multilingualhierarchicalclassificationjob,
  title={Multilingual hierarchical classification of job advertisements for job vacancy statistics},
  author={Maciej Beręsewicz and Marek Wydmuch and Herman Cherniaiev and Robert Pater},
  year={2024},
  eprint={2411.03779},
  archivePrefix={arXiv},
  primaryClass={stat.AP},
  url={https://arxiv.org/abs/2411.03779},
}
```

```
@data{OCUTSI_2024,
  author = {Beręsewicz, Maciej and Wydmuch, Marek and Pater, Robert and Cherniaiev, Herman},
  publisher = {RepOD},
  title = "{Job offers classifiers for ISCO and KZiS 2023}",
  year = {2024},
  version = {V1},
  doi = {10.18150/OCUTSI},
  url = {https://doi.org/10.18150/OCUTSI}
}
```

0.2 Setup (installation and models)

Note that we check the code for the Python versions: 3.9, 3.10 and 3.11.

```
[ ]: !python --version
```

Python 3.11.11

```
[ ]: !rm -rf job_offers_classifier
```

```
[ ]: !git clone https://github.com/OJALAB/job-ads-classifier.git
!mv job-ads-classifier/* ./
!rm -rf job-ads-classifier
```

```
Cloning into 'job-ads-classifier'...
remote: Enumerating objects: 30, done.
remote: Counting objects: 100% (30/30), done.
remote: Compressing objects: 100% (25/25), done.
remote: Total 30 (delta 5), reused 23 (delta 3), pack-reused 0 (from 0)
Receiving objects: 100% (30/30), 2.14 MiB | 13.47 MiB/s, done.
Resolving deltas: 100% (5/5), done.
```

```
[ ]: !pip install virtualenv
!virtualenv classifier
!classifier/bin/pip install --upgrade pip
!classifier/bin/pip install torch
!classifier/bin/pip install -r requirements-colab.txt
```

```
Collecting virtualenv
  Downloading virtualenv-20.29.3-py3-none-any.whl.metadata (4.5 kB)
Collecting distlib<1,>=0.3.7 (from virtualenv)
  Downloading distlib-0.3.9-py2.py3-none-any.whl.metadata (5.2 kB)
Requirement already satisfied: filelock<4,>=3.12.2 in
/usr/local/lib/python3.11/dist-packages (from virtualenv) (3.17.0)
Requirement already satisfied: platformdirs<5,>=3.9.1 in
/usr/local/lib/python3.11/dist-packages (from virtualenv) (4.3.6)
Downloading virtualenv-20.29.3-py3-none-any.whl (4.3 MB)
4.3/4.3 MB
78.5 MB/s eta 0:00:00
Downloading distlib-0.3.9-py2.py3-none-any.whl (468 kB)
469.0/469.0 kB
30.2 MB/s eta 0:00:00
Installing collected packages: distlib, virtualenv
Successfully installed distlib-0.3.9 virtualenv-20.29.3
created virtual environment CPython3.11.11.final.0-64 in 1572ms
  creator CPython3Posix(dest=/content/classifier, clear=False,
no_vcs_ignore=False, global=False)
  seeder FromAppData(download=False, pip=bundle, setuptools=bundle,
wheel=bundle, via=copy, app_data_dir=/root/.local/share/virtualenv)
    added seed packages: pip==25.0.1, setuptools==75.8.0, wheel==0.45.1
  activators BashActivator,CShellActivator,FishActivator,NushellActivator,PowerS
hellActivator,PythonActivator
Requirement already satisfied: pip in ./classifier/lib/python3.11/site-packages
(25.0.1)
Collecting torch
```

Downloading torch-2.6.0-cp311-cp311-manylinux1_x86_64.whl.metadata (28 kB)
Collecting filelock (from torch)
Downloading filelock-3.18.0-py3-none-any.whl.metadata (2.9 kB)
Collecting typing-extensions>=4.10.0 (from torch)
Downloading typing_extensions-4.12.2-py3-none-any.whl.metadata (3.0 kB)
Collecting networkx (from torch)
Downloading networkx-3.4.2-py3-none-any.whl.metadata (6.3 kB)
Collecting jinja2 (from torch)
Downloading jinja2-3.1.6-py3-none-any.whl.metadata (2.9 kB)
Collecting fsspec (from torch)
Downloading fsspec-2025.3.0-py3-none-any.whl.metadata (11 kB)
Collecting nvidia-cuda-nvrtc-cu12==12.4.127 (from torch)
Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cuda-runtime-cu12==12.4.127 (from torch)
Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cuda-cupti-cu12==12.4.127 (from torch)
Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cudnn-cu12==9.1.0.70 (from torch)
Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cublas-cu12==12.4.5.8 (from torch)
Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cufft-cu12==11.2.1.3 (from torch)
Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-curand-cu12==10.3.5.147 (from torch)
Downloading nvidia_curand_cu12-10.3.5.147-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cusolver-cu12==11.6.1.9 (from torch)
Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cusparse-cu12==12.3.1.170 (from torch)
Downloading nvidia_cusparse_cu12-12.3.1.170-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cusparselt-cu12==0.6.2 (from torch)
Downloading nvidia_cusparselt_cu12-0.6.2-py3-none-manylinux2014_x86_64.whl.metadata (6.8 kB)
Collecting nvidia-nccl-cu12==2.21.5 (from torch)
Downloading nvidia_nccl_cu12-2.21.5-py3-none-manylinux2014_x86_64.whl.metadata (1.8 kB)
Collecting nvidia-nvtx-cu12==12.4.127 (from torch)
Downloading nvidia_nvtx_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.7 kB)
Collecting nvidia-nvjitlink-cu12==12.4.127 (from torch)

```

    Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting triton==3.2.0 (from torch)
    Downloading
triton-3.2.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata
(1.4 kB)
Collecting sympy==1.13.1 (from torch)
    Downloading sympy-1.13.1-py3-none-any.whl.metadata (12 kB)
Collecting mpmath<1.4,>=1.1.0 (from sympy==1.13.1->torch)
    Downloading mpmath-1.3.0-py3-none-any.whl.metadata (8.6 kB)
Collecting MarkupSafe>=2.0 (from jinja2->torch)
    Downloading MarkupSafe-3.0.2-cp311-cp311-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (4.0 kB)
Downloading torch-2.6.0-cp311-cp311-manylinux1_x86_64.whl (766.7 MB)
766.7/766.7 MB
29.6 MB/s eta 0:00:00
Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl
(363.4 MB)
363.4/363.4 MB
59.0 MB/s eta 0:00:00
Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl (13.8 MB)
13.8/13.8 MB
109.4 MB/s eta 0:00:00
Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl (24.6 MB)
24.6/24.6 MB
110.1 MB/s eta 0:00:00
Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl (883 kB)
883.7/883.7 kB
48.4 MB/s eta 0:00:00
Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl
(664.8 MB)
664.8/664.8 MB
40.7 MB/s eta 0:00:00
Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl
(211.5 MB)
211.5/211.5 MB
65.6 MB/s eta 0:00:00
Downloading nvidia_curand_cu12-10.3.5.147-py3-none-
manylinux2014_x86_64.whl (56.3 MB)
56.3/56.3 MB
79.2 MB/s eta 0:00:00
Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-
manylinux2014_x86_64.whl (127.9 MB)
127.9/127.9 MB
65.9 MB/s eta 0:00:00

```

Downloading nvidia_cusparselt_cu12-12.3.1.170-py3-none-manylinux2014_x86_64.whl (207.5 MB)
 207.5/207.5 MB
 73.0 MB/s eta 0:00:00
 Downloading nvidia_cusparselt_cu12-0.6.2-py3-none-manylinux2014_x86_64.whl (150.1 MB)
 150.1/150.1 MB
 68.3 MB/s eta 0:00:00
 Downloading nvidia_nccl_cu12-2.21.5-py3-none-manylinux2014_x86_64.whl (188.7 MB)
 188.7/188.7 MB
 66.0 MB/s eta 0:00:00
 Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (21.1 MB)
 21.1/21.1 MB
 110.1 MB/s eta 0:00:00
 Downloading nvidia_nvtx_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (99 kB)
 Downloading sympy-1.13.1-py3-none-any.whl (6.2 MB)
 6.2/6.2 MB
 104.6 MB/s eta 0:00:00
 Downloading triton-3.2.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (253.2 MB)
 253.2/253.2 MB
 54.0 MB/s eta 0:00:00
 Downloading typing_extensions-4.12.2-py3-none-any.whl (37 kB)
 Downloading filelock-3.18.0-py3-none-any.whl (16 kB)
 Downloading fsspec-2025.3.0-py3-none-any.whl (193 kB)
 Downloading jinja2-3.1.6-py3-none-any.whl (134 kB)
 Downloading networkx-3.4.2-py3-none-any.whl (1.7 MB)
 1.7/1.7 MB
 81.4 MB/s eta 0:00:00
 Downloading MarkupSafe-3.0.2-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (23 kB)
 Downloading mpmath-1.3.0-py3-none-any.whl (536 kB)
 536.2/536.2 kB
 27.8 MB/s eta 0:00:00
 Installing collected packages: triton, nvidia-cusparselt-cu12, mpmath, typing-extensions, sympy, nvidia-nvtx-cu12, nvidia-nvjitlink-cu12, nvidia-nccl-cu12, nvidia-curand-cu12, nvidia-cufft-cu12, nvidia-cuda-runtime-cu12, nvidia-cuda-nvrtc-cu12, nvidia-cuda-cupti-cu12, nvidia-cublas-cu12, networkx, MarkupSafe, fsspec, filelock, nvidia-cusparselt-cu12, nvidia-cudnn-cu12, jinja2, nvidia-cusolver-cu12, torch
 Successfully installed MarkupSafe-3.0.2 filelock-3.18.0 fsspec-2025.3.0 jinja2-3.1.6 mpmath-1.3.0 networkx-3.4.2 nvidia-cublas-cu12-12.4.5.8 nvidia-cuda-cupti-cu12-12.4.127 nvidia-cuda-nvrtc-cu12-12.4.127 nvidia-cuda-runtime-

```

cu12-12.4.127 nvidia-cudnn-cu12-9.1.0.70 nvidia-cufft-cu12-11.2.1.3 nvidia-
curand-cu12-10.3.5.147 nvidia-cusolver-cu12-11.6.1.9 nvidia-cusparse-
cu12-12.3.1.170 nvidia-cusparselt-cu12-0.6.2 nvidia-nccl-cu12-2.21.5 nvidia-
nvjitlink-cu12-12.4.127 nvidia-nvtx-cu12-12.4.127 sympy-1.13.1 torch-2.6.0
triton-3.2.0 typing-extensions-4.12.2
Collecting click==8.0.3 (from -r requirements-colab.txt (line 1))
  Downloading click-8.0.3-py3-none-any.whl.metadata (3.2 kB)
Collecting jsbeautifier==1.15.1 (from -r requirements-colab.txt (line 2))
  Downloading jsbeautifier-1.15.1.tar.gz (75 kB)
  Installing build dependencies ... done
  Getting requirements to build wheel ... done
  Preparing metadata (pyproject.toml) ... done
Collecting lxml==5.3.0 (from -r requirements-colab.txt (line 3))
  Downloading lxml-5.3.0-cp311-cp311-manylinux_2_28_x86_64.whl.metadata (3.8 kB)
Collecting napkinxc==0.7.1 (from -r requirements-colab.txt (line 4))
  Downloading napkinxc-0.7.1-cp311-cp311-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (6.0 kB)
Collecting numpy==1.24.3 (from -r requirements-colab.txt (line 5))
  Downloading
numpy-1.24.3-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata
(5.6 kB)
Collecting pandas==2.2.2 (from -r requirements-colab.txt (line 6))
  Downloading
pandas-2.2.2-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata
(19 kB)
Collecting pyreadr==0.5.2 (from -r requirements-colab.txt (line 7))
  Downloading pyreadr-0.5.2-cp311-cp311-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (1.2 kB)
Collecting pystempel==1.2.0 (from -r requirements-colab.txt (line 8))
  Downloading pystempel-1.2.0-py3-none-any.whl.metadata (7.9 kB)
Collecting pytorch_lightning>=2.0.0 (from -r requirements-colab.txt (line 9))
  Downloading pytorch_lightning-2.5.0.post0-py3-none-any.whl.metadata (21 kB)
Collecting sacremoses==0.1.1 (from -r requirements-colab.txt (line 10))
  Downloading sacremoses-0.1.1-py3-none-any.whl.metadata (8.3 kB)
Collecting scikit_learn==1.2.2 (from -r requirements-colab.txt (line 11))
  Downloading scikit_learn-1.2.2-cp311-cp311-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (11 kB)
Collecting scipy>=1.11.0 (from -r requirements-colab.txt (line 12))
  Downloading
scipy-1.15.2-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata
(61 kB)
Collecting stop_words==2018.7.23 (from -r requirements-colab.txt (line 13))
  Downloading stop-words-2018.7.23.tar.gz (31 kB)
  Preparing metadata (setup.py) ... done
Collecting torchmetrics==1.4.1 (from -r requirements-colab.txt (line 14))
  Downloading torchmetrics-1.4.1-py3-none-any.whl.metadata (20 kB)
Collecting tqdm==4.66.5 (from -r requirements-colab.txt (line 15))
  Downloading tqdm-4.66.5-py3-none-any.whl.metadata (57 kB)

```

Collecting transformers==4.44.2 (from -r requirements-colab.txt (line 16))
 Downloading transformers-4.44.2-py3-none-any.whl.metadata (43 kB)
 Collecting six>=1.13.0 (from jsbeautifier==1.15.1->-r requirements-colab.txt (line 2))
 Downloading six-1.17.0-py2.py3-none-any.whl.metadata (1.7 kB)
 Collecting editorconfig>=0.12.2 (from jsbeautifier==1.15.1->-r requirements-colab.txt (line 2))
 Downloading EditorConfig-0.17.0-py3-none-any.whl.metadata (3.8 kB)
 Collecting gdown (from napkinxc==0.7.1->-r requirements-colab.txt (line 4))
 Downloading gdown-5.2.0-py3-none-any.whl.metadata (5.8 kB)
 Collecting python-dateutil>=2.8.2 (from pandas==2.2.2->-r requirements-colab.txt (line 6))
 Downloading python_dateutil-2.9.0.post0-py2.py3-none-any.whl.metadata (8.4 kB)
 Collecting pytz>=2020.1 (from pandas==2.2.2->-r requirements-colab.txt (line 6))
 Downloading pytz-2025.1-py2.py3-none-any.whl.metadata (22 kB)
 Collecting tzdata>=2022.7 (from pandas==2.2.2->-r requirements-colab.txt (line 6))
 Downloading tzdata-2025.1-py2.py3-none-any.whl.metadata (1.4 kB)
 Collecting sortedcontainers (from pystempel==1.2.0->-r requirements-colab.txt (line 8))
 Downloading sortedcontainers-2.4.0-py2.py3-none-any.whl.metadata (10 kB)
 Collecting regex (from sacremoses==0.1.1->-r requirements-colab.txt (line 10))
 Downloading regex-2024.11.6-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (40 kB)
 Collecting joblib (from sacremoses==0.1.1->-r requirements-colab.txt (line 10))
 Downloading joblib-1.4.2-py3-none-any.whl.metadata (5.4 kB)
 Collecting threadpoolctl>=2.0.0 (from scikit_learn==1.2.2->-r requirements-colab.txt (line 11))
 Downloading threadpoolctl-3.6.0-py3-none-any.whl.metadata (13 kB)
 Collecting packaging>17.1 (from torchmetrics==1.4.1->-r requirements-colab.txt (line 14))
 Downloading packaging-24.2-py3-none-any.whl.metadata (3.2 kB)
 Requirement already satisfied: torch>=1.10.0 in ./classifier/lib/python3.11/site-packages (from torchmetrics==1.4.1->-r requirements-colab.txt (line 14)) (2.6.0)
 Collecting lightning-utilities>=0.8.0 (from torchmetrics==1.4.1->-r requirements-colab.txt (line 14))
 Downloading lightning_utilities-0.14.1-py3-none-any.whl.metadata (5.6 kB)
 Requirement already satisfied: filelock in ./classifier/lib/python3.11/site-packages (from transformers==4.44.2->-r requirements-colab.txt (line 16)) (3.18.0)
 Collecting huggingface-hub<1.0,>=0.23.2 (from transformers==4.44.2->-r requirements-colab.txt (line 16))
 Downloading huggingface_hub-0.29.3-py3-none-any.whl.metadata (13 kB)
 Collecting pyyaml>=5.1 (from transformers==4.44.2->-r requirements-colab.txt (line 16))
 Downloading PyYAML-6.0.2-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata

(2.1 kB)
Collecting requests (from transformers==4.44.2->-r requirements-colab.txt (line 16))
Downloading requests-2.32.3-py3-none-any.whl.metadata (4.6 kB)
Collecting safetensors>=0.4.1 (from transformers==4.44.2->-r requirements-colab.txt (line 16))
Downloading safetensors-0.5.3-cp38-abi3-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (3.8 kB)
Collecting tokenizers<0.20,>=0.19 (from transformers==4.44.2->-r requirements-colab.txt (line 16))
Downloading tokenizers-0.19.1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (6.7 kB)
Requirement already satisfied: fsspec>=2022.5.0 in ./classifier/lib/python3.11/site-packages (from fsspec[http]>=2022.5.0->pytorch_lightning>=2.0.0->-r requirements-colab.txt (line 9)) (2025.3.0)
Requirement already satisfied: typing-extensions>=4.4.0 in ./classifier/lib/python3.11/site-packages (from pytorch_lightning>=2.0.0->-r requirements-colab.txt (line 9)) (4.12.2)
Collecting aiohttp!=4.0.0a0,!4.0.0a1 (from fsspec[http]>=2022.5.0->pytorch_lightning>=2.0.0->-r requirements-colab.txt (line 9))
Downloading aiohttp-3.11.14-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (7.7 kB)
Requirement already satisfied: setuptools in ./classifier/lib/python3.11/site-packages (from lightning-utilities>=0.8.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14)) (75.8.0)
Requirement already satisfied: networkx in ./classifier/lib/python3.11/site-packages (from torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14)) (3.4.2)
Requirement already satisfied: jinja2 in ./classifier/lib/python3.11/site-packages (from torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14)) (3.1.6)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.4.127 in ./classifier/lib/python3.11/site-packages (from torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14)) (12.4.127)
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.4.127 in ./classifier/lib/python3.11/site-packages (from torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14)) (12.4.127)
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.4.127 in ./classifier/lib/python3.11/site-packages (from torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14)) (12.4.127)
Requirement already satisfied: nvidia-cudnn-cu12==9.1.0.70 in ./classifier/lib/python3.11/site-packages (from torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14))


```

(9.1.0.70)
Requirement already satisfied: nvidia-cublas-cu12==12.4.5.8 in
./classifier/lib/python3.11/site-packages (from
torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14))
(12.4.5.8)
Requirement already satisfied: nvidia-cufft-cu12==11.2.1.3 in
./classifier/lib/python3.11/site-packages (from
torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14))
(11.2.1.3)
Requirement already satisfied: nvidia-curand-cu12==10.3.5.147 in
./classifier/lib/python3.11/site-packages (from
torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14))
(10.3.5.147)
Requirement already satisfied: nvidia-cusolver-cu12==11.6.1.9 in
./classifier/lib/python3.11/site-packages (from
torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14))
(11.6.1.9)
Requirement already satisfied: nvidia-cusparse-cu12==12.3.1.170 in
./classifier/lib/python3.11/site-packages (from
torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14))
(12.3.1.170)
Requirement already satisfied: nvidia-cusparselt-cu12==0.6.2 in
./classifier/lib/python3.11/site-packages (from
torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14)) (0.6.2)
Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in
./classifier/lib/python3.11/site-packages (from
torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14))
(2.21.5)
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in
./classifier/lib/python3.11/site-packages (from
torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14))
(12.4.127)
Requirement already satisfied: nvidia-nvjitlink-cu12==12.4.127 in
./classifier/lib/python3.11/site-packages (from
torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14))
(12.4.127)
Requirement already satisfied: triton==3.2.0 in
./classifier/lib/python3.11/site-packages (from
torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14)) (3.2.0)
Requirement already satisfied: sympy==1.13.1 in
./classifier/lib/python3.11/site-packages (from
torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14))
(1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in
./classifier/lib/python3.11/site-packages (from
sympy==1.13.1->torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt
(line 14)) (1.3.0)
Collecting beautifulsoup4 (from gdown->napkinxc==0.7.1->-r requirements-

```

```

colab.txt (line 4))
  Downloading beautifulsoup4-4.13.3-py3-none-any.whl.metadata (3.8 kB)
Collecting charset-normalizer<4,>=2 (from requests->transformers==4.44.2->-r
requirements-colab.txt (line 16))
  Downloading charset_normalizer-3.4.1-cp311-cp311-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (35 kB)
Collecting idna<4,>=2.5 (from requests->transformers==4.44.2->-r requirements-
colab.txt (line 16))
  Downloading idna-3.10-py3-none-any.whl.metadata (10 kB)
Collecting urllib3<3,>=1.21.1 (from requests->transformers==4.44.2->-r
requirements-colab.txt (line 16))
  Downloading urllib3-2.3.0-py3-none-any.whl.metadata (6.5 kB)
Collecting certifi>=2017.4.17 (from requests->transformers==4.44.2->-r
requirements-colab.txt (line 16))
  Downloading certifi-2025.1.31-py3-none-any.whl.metadata (2.5 kB)
Collecting aiohappyeyeballs>=2.3.0 (from
aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]>=2022.5.0->pytorch_lightning>=2.0.0->-r
requirements-colab.txt (line 9))
  Downloading aiohappyeyeballs-2.6.1-py3-none-any.whl.metadata (5.9 kB)
Collecting aiosignal>=1.1.2 (from
aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]>=2022.5.0->pytorch_lightning>=2.0.0->-r
requirements-colab.txt (line 9))
  Downloading aiosignal-1.3.2-py2.py3-none-any.whl.metadata (3.8 kB)
Collecting attrs>=17.3.0 (from
aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]>=2022.5.0->pytorch_lightning>=2.0.0->-r
requirements-colab.txt (line 9))
  Downloading attrs-25.3.0-py3-none-any.whl.metadata (10 kB)
Collecting frozenlist>=1.1.1 (from
aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]>=2022.5.0->pytorch_lightning>=2.0.0->-r
requirements-colab.txt (line 9))
  Downloading frozenlist-1.5.0-cp311-cp311-
manylinux_2_5_x86_64.manylinux1_x86_64.manylinux_2_17_x86_64.manylinux2014_x86_6
4.whl.metadata (13 kB)
Collecting multidict<7.0,>=4.5 (from
aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]>=2022.5.0->pytorch_lightning>=2.0.0->-r
requirements-colab.txt (line 9))
  Downloading multidict-6.2.0-cp311-cp311-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (4.9 kB)
Collecting propcache>=0.2.0 (from
aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]>=2022.5.0->pytorch_lightning>=2.0.0->-r
requirements-colab.txt (line 9))
  Downloading propcache-0.3.0-cp311-cp311-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (10 kB)
Collecting yarl<2.0,>=1.17.0 (from
aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]>=2022.5.0->pytorch_lightning>=2.0.0->-r
requirements-colab.txt (line 9))
  Downloading
yarl-1.18.3-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata

```

```

(69 kB)
Collecting soupsieve>1.2 (from beautifulsoup4->gdown->napkinxc==0.7.1->-r
requirements-colab.txt (line 4))
  Downloading soupsieve-2.6-py3-none-any.whl.metadata (4.6 kB)
Requirement already satisfied: MarkupSafe>=2.0 in
./classifier/lib/python3.11/site-packages (from
jinja2->torch>=1.10.0->torchmetrics==1.4.1->-r requirements-colab.txt (line 14))
(3.0.2)
Collecting PySocks!=1.5.7,>=1.5.6 (from
requests[socks]->gdown->napkinxc==0.7.1->-r requirements-colab.txt (line 4))
  Downloading PySocks-1.7.1-py3-none-any.whl.metadata (13 kB)
Downloading click-8.0.3-py3-none-any.whl (97 kB)
Downloading lxml-5.3.0-cp311-cp311-manylinux_2_28_x86_64.whl (5.0 MB)
      5.0/5.0 MB
118.8 MB/s eta 0:00:00
Downloading
napkinxc-0.7.1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (555
kB)
      555.1/555.1 kB
28.0 MB/s eta 0:00:00
Downloading
numpy-1.24.3-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (17.3
MB)
      17.3/17.3 MB
122.1 MB/s eta 0:00:00
Downloading
pandas-2.2.2-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (13.0
MB)
      13.0/13.0 MB
125.5 MB/s eta 0:00:00
Downloading
pyreadr-0.5.2-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (416
kB)
Downloading pystempel-1.2.0-py3-none-any.whl (2.7 MB)
      2.7/2.7 MB
9.2 MB/s eta 0:00:00
Downloading sacremoses-0.1.1-py3-none-any.whl (897 kB)
      897.5/897.5 kB
46.6 MB/s eta 0:00:00
Downloading
scikit_learn-1.2.2-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl
(9.6 MB)
      9.6/9.6 MB
133.9 MB/s eta 0:00:00
Downloading torchmetrics-1.4.1-py3-none-any.whl (866 kB)
      866.2/866.2 kB
43.5 MB/s eta 0:00:00
Downloading tqdm-4.66.5-py3-none-any.whl (78 kB)

```

```

Downloading transformers-4.44.2-py3-none-any.whl (9.5 MB)
9.5/9.5 MB
133.2 MB/s eta 0:00:00
Downloading pytorch_lightning-2.5.0.post0-py3-none-any.whl (819 kB)
819.3/819.3 kB
43.4 MB/s eta 0:00:00
Downloading
scipy-1.15.2-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (37.6 MB)
37.6/37.6 MB
51.2 MB/s eta 0:00:00
Downloading EditorConfig-0.17.0-py3-none-any.whl (16 kB)
Downloading huggingface_hub-0.29.3-py3-none-any.whl (468 kB)
Downloading joblib-1.4.2-py3-none-any.whl (301 kB)
Downloading lightning_utilities-0.14.1-py3-none-any.whl (28 kB)
Downloading packaging-24.2-py3-none-any.whl (65 kB)
Downloading python_dateutil-2.9.0.post0-py2.py3-none-any.whl (229 kB)
Downloading pytz-2025.1-py2.py3-none-any.whl (507 kB)
Downloading
PyYAML-6.0.2-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (762 kB)
763.0/763.0 kB
41.8 MB/s eta 0:00:00
Downloading
regex-2024.11.6-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (792 kB)
792.7/792.7 kB
47.3 MB/s eta 0:00:00
Downloading
safetensors-0.5.3-cp38-abi3-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (471 kB)
Downloading six-1.17.0-py2.py3-none-any.whl (11 kB)
Downloading threadpoolctl-3.6.0-py3-none-any.whl (18 kB)
Downloading
tokenizers-0.19.1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (3.6 MB)
3.6/3.6 MB
126.7 MB/s eta 0:00:00
Downloading tzdata-2025.1-py2.py3-none-any.whl (346 kB)
Downloading gdown-5.2.0-py3-none-any.whl (18 kB)
Downloading requests-2.32.3-py3-none-any.whl (64 kB)
Downloading sortedcontainers-2.4.0-py2.py3-none-any.whl (29 kB)
Downloading
aiohttp-3.11.14-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (1.7 MB)
1.7/1.7 MB
87.4 MB/s eta 0:00:00
Downloading certifi-2025.1.31-py3-none-any.whl (166 kB)
Downloading charset_normalizer-3.4.1-cp311-cp311-

```

```

manylinux_2_17_x86_64.manylinux2014_x86_64.whl (143 kB)
Downloading idna-3.10-py3-none-any.whl (70 kB)
Downloading urllib3-2.3.0-py3-none-any.whl (128 kB)
Downloading beautifulsoup4-4.13.3-py3-none-any.whl (186 kB)
Downloading aiohappyeyeballs-2.6.1-py3-none-any.whl (15 kB)
Downloading aiosignal-1.3.2-py2.py3-none-any.whl (7.6 kB)
Downloading attrs-25.3.0-py3-none-any.whl (63 kB)
Downloading frozenlist-1.5.0-cp311-cp311-
manylinux_2_5_x86_64.manylinux1_x86_64.manylinux_2_17_x86_64.manylinux2014_x86_6
4.whl (274 kB)
Downloading
multidict-6.2.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (133
kB)
Downloading
propcache-0.3.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (231
kB)
Downloading PySocks-1.7.1-py3-none-any.whl (16 kB)
Downloading soupsieve-2.6-py3-none-any.whl (36 kB)
Downloading
yarl-1.18.3-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (344 kB)
Building wheels for collected packages: jsbeautifier, stop_words
  Building wheel for jsbeautifier (pyproject.toml) ... done
  Created wheel for jsbeautifier: filename=jsbeautifier-1.15.1-py3-none-any.whl
size=94751
sha256=d6069099846fb80b388cb5b892362e4a9d444483290d3b749cf9a406b0b3865f
  Stored in directory: /root/.cache/pip/wheels/32/eb/d8/07a66b0cf535ed05084d0b9a
920b89e35f20512a5e28e615ca
  Building wheel for stop_words (setup.py) ... done
  Created wheel for stop_words: filename=stop_words-2018.7.23-py3-none-any.whl
size=32940
sha256=d73f5185f12c98302f7a04d85963471462e3c54fe5a20d95553d7bcdece5f16e
  Stored in directory: /root/.cache/pip/wheels/8f/a5/51/a5405e1da5d178491b79d12c
c81b6cb9bb14fe2c8c632eba70
Successfully built jsbeautifier stop_words
Installing collected packages: stop_words, sortedcontainers, pytz, editorconfig,
urllib3, tzdata, tqdm, threadpoolctl, soupsieve, six, safetensors, regex,
pyyaml, PySocks, propcache, packaging, numpy, multidict, lxml, joblib, idna,
frozenlist, click, charset-normalizer, certifi, attrs, aiohappyeyeballs, yarl,
scipy, sacremoses, requests, python-dateutil, pystempel, lightning-utilities,
jsbeautifier, beautifulsoup4, aiosignal, scikit_learn, pandas, huggingface-hub,
aiohttp, torchmetrics, tokenizers, pyreadr, gdown, transformers,
pytorch_lightning, napkinxc
Successfully installed PySocks-1.7.1 aiohappyeyeballs-2.6.1 aiohttp-3.11.14
aiosignal-1.3.2 attrs-25.3.0 beautifulsoup4-4.13.3 certifi-2025.1.31 charset-
normalizer-3.4.1 click-8.0.3 editorconfig-0.17.0 frozenlist-1.5.0 gdown-5.2.0
huggingface-hub-0.29.3 idna-3.10 joblib-1.4.2 jsbeautifier-1.15.1 lightning-
utilities-0.14.1 lxml-5.3.0 multidict-6.2.0 napkinxc-0.7.1 numpy-1.24.3
packaging-24.2 pandas-2.2.2 propcache-0.3.0 pyreadr-0.5.2 pystempel-1.2.0

```

```
python-dateutil-2.9.0.post0 pytorch_lightning-2.5.0.post0 pytz-2025.1
pyyaml-6.0.2 regex-2024.11.6 requests-2.32.3 sacremoses-0.1.1 safetensors-0.5.3
scikit_learn-1.2.2 scipy-1.15.2 six-1.17.0 sortedcontainers-2.4.0 soupsieve-2.6
stop_words-2018.7.23 threadpoolctl-3.6.0 tokenizers-0.19.1 torchmetrics-1.4.1
tqdm-4.66.5 transformers-4.44.2 tzdata-2025.1 urllib3-2.3.0 yar1-1.18.3
```

```
[ ]: import sys
      sys.path.append("classifier/lib/python3.11/site-packages")
```

Download multilingual models (around 20 min for each file)

- top-down (<https://repod.icm.edu.pl/api/access/datafile/50892>)
- bottom-up (<https://repod.icm.edu.pl/api/access/datafile/50893>)

```
[ ]: !wget https://repod.icm.edu.pl/api/access/datafile/50892 -O top-down.zip
```

```
--2025-03-18 09:38:11-- https://repod.icm.edu.pl/api/access/datafile/50892
Resolving repod.icm.edu.pl (repod.icm.edu.pl)... 213.135.60.199
Connecting to repod.icm.edu.pl (repod.icm.edu.pl)|213.135.60.199|:443...
connected.
HTTP request sent, awaiting response... 200 OK
Length: 2117465081 (2.0G) [application/zip]
Saving to: 'top-down.zip'
```

```
top-down.zip      100%[=====>]    1.97G  4.24MB/s   in 15m 17s
```

```
2025-03-18 09:53:30 (2.20 MB/s) - 'top-down.zip' saved [2117465081/2117465081]
```

```
[ ]: !wget https://repod.icm.edu.pl/api/access/datafile/50893 -O bottom-up.zip
```

```
--2025-03-18 09:54:44-- https://repod.icm.edu.pl/api/access/datafile/50893
Resolving repod.icm.edu.pl (repod.icm.edu.pl)... 213.135.60.199
Connecting to repod.icm.edu.pl (repod.icm.edu.pl)|213.135.60.199|:443...
connected.
HTTP request sent, awaiting response... 200 OK
Length: 3049932246 (2.8G) [application/zip]
Saving to: 'bottom-up.zip'
```

```
bottom-up.zip     100%[=====>]    2.84G  3.31MB/s   in 16m 5s
```

```
2025-03-18 10:10:50 (3.01 MB/s) - 'bottom-up.zip' saved [3049932246/3049932246]
```

```
[ ]: !unzip -l top-down.zip
```

```
Archive:  top-down.zip
  Length   Date      Time     Name
  -----  -
           0  2024-10-23 10:26  transformer-multi-top-2024-sample-20/
```

```

0 2024-10-25 06:48 transformer-multi-top-2024-sample-20/ckpts/
3392393018 2024-10-24 02:17 transformer-multi-
top-2024-sample-20/ckpts/epoch=4-val_loss=0.31984.ckpt
1319109 2024-10-23 06:26 transformer-multi-top-2024-sample-20/hierarchy.bin
88 2024-10-23 06:26 transformer-multi-
top-2024-sample-20/transformer_arch.bin
-----
3392712215 5 files

```

```
[ ]: !unzip top-down.zip
```

```

Archive: top-down.zip
  creating: transformer-multi-top-2024-sample-20/
  creating: transformer-multi-top-2024-sample-20/ckpts/
  inflating: transformer-multi-
top-2024-sample-20/ckpts/epoch=4-val_loss=0.31984.ckpt
  inflating: transformer-multi-top-2024-sample-20/hierarchy.bin
  inflating: transformer-multi-top-2024-sample-20/transformer_arch.bin

```

```
[ ]: !unzip bottom-up.zip
```

```

Archive: bottom-up.zip
  creating: transformer-multi-bottom-2024-sample-20/
  inflating: transformer-multi-bottom-2024-sample-20/hierarchy.bin
  inflating: transformer-multi-bottom-2024-sample-20/transformer_arch.bin
  creating: transformer-multi-bottom-2024-sample-20/ckpts/
  inflating: transformer-multi-
bottom-2024-sample-20/ckpts/epoch=6-val_loss=1.59784.ckpt
  inflating: transformer-multi-bottom-2024-sample-20/config.json
  inflating: transformer-multi-bottom-2024-sample-20/model.safetensors

```

0.3 Loading library and simple prediction

Libraries and functions

```
[ ]: import numpy as np
import pandas as pd
from job_offers_classifier.load_save import load_to_df, save_obj, load_obj ##
    ↳ here it is not needed actually
from job_offers_classifier.job_offers_utils import create_hierarchy,
    ↳ fix_class_str, remove_classes, filter_classes, remap_classes ## here it is
    ↳ not needed actually
from job_offers_classifier.job_offers_classifier import *
```

Read the models

```
[ ]: trans_bottom = TransformerJobOffersClassifier(modeling_mode="bottom-up")
trans_bottom.load("transformer-multi-bottom-2024-sample-20")
```

```
trans_top = TransformerJobOffersClassifier(modeling_mode="top-down")
trans_top.load("transformer-multi-top-2024-sample-20")
```

```
Initializing TransformerClassifier with model_name=FacebookAI/xlm-roberta-base,
  output_size=2911, labels_groups=False,
  learning_rate=1e-05, weight_decay=0.01, warmup_steps=50, train_batch_size=8,
eval_batch_size=8, freeze_transformer=False
  loss=flat cross entropy, hidden_dropout=0.0, hierarchy_leaves=2911 ...
```

```
/usr/local/lib/python3.11/dist-packages/huggingface_hub/utils/_auth.py:94:
```

```
UserWarning:
```

```
The secret `HF_TOKEN` does not exist in your Colab secrets.
```

```
To authenticate with the Hugging Face Hub, create a token in your settings tab
(https://huggingface.co/settings/tokens), set it as secret in your Google Colab
and restart your session.
```

```
You will be able to reuse this secret in all of your notebooks.
```

```
Please note that authentication is recommended but still optional to access
public models or datasets.
```

```
warnings.warn(
```

```
config.json: 0%|          | 0.00/615 [00:00<?, ?B/s]
```

```
model.safetensors: 0%|          | 0.00/1.12G [00:00<?, ?B/s]
```

```
Initializing TransformerClassifier with model_name=FacebookAI/xlm-roberta-base,
  output_size=3543, labels_groups=True,
  learning_rate=1e-05, weight_decay=0.01, warmup_steps=50, train_batch_size=8,
eval_batch_size=8, freeze_transformer=False
  loss=hierarchical cross entropy, hidden_dropout=0.0, hierarchy_leaves=2911 ...
```

```
Example offers
```

```
[ ]: example_offer = [
```


"Starszy Statystyk w Ośrodku Metodologii Badań Ludnościowych OSOBA NA TYM

- STANOWISKU: Bierze współudział w rozwoju metodologii estymacji wielkości
- populacji i jej charakterystyk na podstawie zintegrowanych danych Prowadzi
- podstawowe prace projektowo-programistyczne dotyczące estymacji i integracji
- danych Współpracuje z departamentami GUS, urzędami statystycznymi i
- komórkami organizacyjnymi Urzędu w zakresie estymacji wielkości populacji i
- ich charakterystyk na podstawie zintegrowanych danych Uczestniczy we
- współpracy międzynarodowej w dziedzinie estymacji wielkości populacji i ich
- charakterystyk na podstawie zintegrowanych danych Współpracuje przy
- przygotowywaniu danych niezbędnych do realizacji zleceń zewnętrznych
- związanych z udostępnianiem informacji Bierze udział w rozpoznawaniu
- rejestrów administracyjnych jako baz danych ludnościowych Wykształcenie:
- średnie Znajomość języka angielskiego na poziomie komunikatywnym Podstawowa
- znajomość metod statystycznych ze szczególnym uwzględnieniem metody
- reprezentacyjnej Podstawowa znajomość pakietów statystycznych R lub Julia
- (ewentualnie Python) Umiejętność argumentowania Rzetelność Organizacja pracy
- i orientacja na osiąganie celów Wykorzystywanie wiedzy i doskonalenie
- zawodowe Współpraca Komunikacja interpersonalna Wykształcenie: wyższe
- profilowane (matematyka, informatyka, ekonomia) Doświadczenie zawodowe co
- najmniej 1 rok w obszarze metodologii badań statystycznych Znajomość
- programu badań statystycznych statystyki publicznej Myślenie analityczne
- Kreatywność",

"Specjalista w Ośrodku Metodologii Badań Ludnościowych Koordynuje prace

- związane z projektami realizowanymi w Ośrodku Metodologii Badań
- Ludnościowych Przygotowuje dokumenty związane z realizacją projektów
- Współpracuje z departamentami GUS Wykształcenie: wyższe Znajomość języka
- angielskiego na poziomie komunikatywnym Umiejętność zarządzania projektami
- Umiejętność argumentowania Rzetelność Organizacja pracy i orientacja na
- osiąganie celów Wykorzystywanie wiedzy i doskonalenie zawodowe Współpraca
- Komunikacja pisemna Komunikacja interpersonalna Wykształcenie: wyższe
- profilowane (ekonomia lub zarządzanie) Doświadczenie zawodowe co najmniej 1
- rok w obszarze metodologii badań statystycznych Znajomość języka
- angielskiego na poziomie bardzo dobrym Znajomość programu badań
- statystycznych statystyki publicznej Myślenie analityczne Kreatywność",

```

    "We are seeking a highly skilled AI Solutions Engineer to join our team.
    ↳The ideal candidate will have extensive experience in machine learning (ML)
    ↳and Generative AI, focusing on both the strategic and tactical development
    ↳of AI solutions. You will architect and deploy advanced AI systems,
    ↳particularly around Generative AI, leveraging your deep technical expertise
    ↳to build innovative products, solve complex problems and guide AI-driven
    ↳strategies for the business. Bachelor's or Master's degree in Computer
    ↳Science, Machine Learning, or a related field (PhD preferred). 5+ years of
    ↳experience in building and deploying machine learning models, with proven
    ↳experience in Generative AI. Expertise in Python and ML frameworks, and deep
    ↳learning architectures for generative models. Proven experience in cloud
    ↳platforms such as AWS or Azure for AI solution deployment. Experience in
    ↳developing products utilizing Generative AI models, such as image
    ↳generation, natural language generation, or synthetic data generation.
    ↳Strong understanding of ML algorithms (supervised, unsupervised,
    ↳reinforcement learning) and Generative AI techniques. Experience in scaling
    ↳generative models and deploying them in production environments. Excellent
    ↳communication skills to articulate technical concepts to non-technical
    ↳stakeholders. Experience in NLP, computer vision, and deploying large
    ↳language models (LLMs) for content generation - is an advantage. Hands-on
    ↳experience with fine-tuning LLMs - is an advantage. Experience with MLOps
    ↳frameworks and version control of models - is nice to have. Knowledge of
    ↳product design principles using AI and ability to drive cross-functional AI
    ↳projects with a generative component - is an asset.",
    "
    /
    /
    ,
    /
    Python,
    ,
    (Gen)AI.
    ,
    -
    ,
    LLM.
    ,
    Python,
    (git,
    ).
    , GenAI
    .
    AI
    .
    ,
    ,
    .
    5
    .
    Python.
    .
    LLM RAG
    (Langchain, LLamaIndex
    ).
    NLP (Natural Language Processing).
    .
    .
    "
]

```

Prediction of the 1-digit occupation code

```

[ ]: pred_result_bottom = trans_bottom.predict(
    example_offer, # text data for which to make a prediction

```

```

    output_level=0, # the layer of the hierarchy for which to make a prediction
    ↪(add a number from 0 - the first level, or 'last' if predicting for the last)
    top_k=None, # how many classes with the highest probability to predict, if
    ↪None is returned the probability for all classes on the selected level of
    ↪the hierarchy
    format='dataframe' # defines the format in which the result will be
    ↪returned, possible options are 'array' - numpy.ndarray or 'dataframe' -
    ↪pandas.DataFrame
)

```

Initializing TransformerDataModule with model_name=FacebookAI/xlm-roberta-base,
max_seq_length=512, train/eval_batch_size=8/8, num_workers=2 ...

Setting up TransformerDataModule ...

tokenizer_config.json: 0%| | 0.00/25.0 [00:00<?, ?B/s]

sentencepiece.bpe.model: 0%| | 0.00/5.07M [00:00<?, ?B/s]

tokenizer.json: 0%| | 0.00/9.10M [00:00<?, ?B/s]

/content/classifier/lib/python3.11/site-

packages/lightning_fabric/connector.py:572: `precision=16` is supported for
historical reasons but its usage is discouraged. Please set your precision to
16-mixed instead!

INFO:pytorch_lightning.utilities.rank_zero:Using 16bit Automatic Mixed Precision
(AMP)

INFO:pytorch_lightning.utilities.rank_zero:GPU available: True (cuda), used:
True

INFO:pytorch_lightning.utilities.rank_zero:TPU available: False, using: 0 TPU
cores

INFO:pytorch_lightning.utilities.rank_zero:HPU available: False, using: 0 HPUs

Starting predicting with TransformerClassifier ...

INFO:pytorch_lightning.utilities.rank_zero:Restoring states from the checkpoint
path at transformer-multi-

bottom-2024-sample-20/ckpts/epoch=6-val_loss=1.59784.ckpt

INFO:pytorch_lightning.utilities.migration.utils:Lightning automatically
upgraded your loaded checkpoint from v1.9.5 to v2.5.0.post0. To apply the
upgrade to your files permanently, run `python -m

pytorch_lightning.utilities.upgrade_checkpoint transformer-multi-

bottom-2024-sample-20/ckpts/epoch=6-val_loss=1.59784.ckpt`

/content/classifier/lib/python3.11/site-

packages/pytorch_lightning/callbacks/model_checkpoint.py:362: The dirpath has
changed from 'models-paper/multi/transformer-multi-bottom-2024-sample-20/ckpts'
to '/content/transformer-multi-

bottom-2024-sample-20/ckpts/epoch=6-val_loss=1.59784.ckpt', therefore

`best_model_score`, `kth_best_model_path`, `kth_value`, `last_model_path` and
`best_k_models` won't be reloaded. Only `best_model_path` will be reloaded.

INFO:pytorch_lightning.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES:

[0]

INFO:pytorch_lightning.utilities.rank_zero:Loaded model weights from the checkpoint at transformer-multi-bottom-2024-sample-20/ckpts/epoch=6-val_loss=1.59784.ckpt

Predicting: | | 0/? [00:00<?, ?it/s]

```
[ ]: pred_result_top = trans_top.predict(  
    example_offer,  
    output_level=0,  
    top_k=None,  
    format="dataframe"  
)
```

Initializing TransformerDataModule with model_name=FacebookAI/xlm-roberta-base, max_seq_length=512, train/eval_batch_size=8/8, num_workers=2 ...
Setting up TransformerDataModule ...

INFO:pytorch_lightning.utilities.rank_zero:Using 16bit Automatic Mixed Precision (AMP)

INFO:pytorch_lightning.utilities.rank_zero:GPU available: True (cuda), used: True

INFO:pytorch_lightning.utilities.rank_zero:TPU available: False, using: 0 TPU cores

INFO:pytorch_lightning.utilities.rank_zero:HPU available: False, using: 0 HPUs

INFO:pytorch_lightning.utilities.rank_zero:Restoring states from the checkpoint path at transformer-multi-top-2024-sample-20/ckpts/epoch=4-val_loss=0.31984.ckpt

Starting predicting with TransformerClassifier ...

INFO:pytorch_lightning.utilities.migration.utils:Lightning automatically upgraded your loaded checkpoint from v1.9.5 to v2.5.0.post0. To apply the upgrade to your files permanently, run `python -m pytorch_lightning.utilities.upgrade_checkpoint transformer-multi-top-2024-sample-20/ckpts/epoch=4-val_loss=0.31984.ckpt`
/content/classifier/lib/python3.11/site-packages/pytorch_lightning/callbacks/model_checkpoint.py:362: The dirpath has changed from 'models-paper/multi/transformer-multi-top-2024-sample-20/ckpts' to '/content/transformer-multi-top-2024-sample-20/ckpts/epoch=4-val_loss=0.31984.ckpt', therefore `best_model_score`, `kth_best_model_path`, `kth_value`, `last_model_path` and `best_k_models` won't be reloaded. Only `best_model_path` will be reloaded.
INFO:pytorch_lightning.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES:

[0]

INFO:pytorch_lightning.utilities.rank_zero:Loaded model weights from the checkpoint at transformer-multi-top-2024-sample-20/ckpts/epoch=4-val_loss=0.31984.ckpt

Predicting: | | 0/? [00:00<?, ?it/s]

Predictions for each 1-digit occupation code for both models (2 two rows: bottop-up model, next

two: top-down model).

```
[ ]: pd.concat([pred_result_bottom, pred_result_top])
```

```
[ ]:
```

	0	1	2	3	4	5	6	\
0	0.000042	0.006040	0.404478	0.237416	0.313389	0.007069	0.001013	
1	0.000012	0.063933	0.609985	0.038708	0.259444	0.008965	0.000519	
2	0.000014	0.004593	0.960370	0.011890	0.002996	0.003023	0.000908	
3	0.000012	0.002445	0.970239	0.012542	0.005935	0.001242	0.000462	
0	0.000312	0.002083	0.828137	0.052120	0.114958	0.000511	0.000613	
1	0.000269	0.123873	0.634011	0.002438	0.236913	0.001435	0.000465	
2	0.000951	0.010495	0.983327	0.000168	0.000621	0.001219	0.001900	
3	0.000598	0.001582	0.992644	0.000637	0.003383	0.000294	0.000585	
	7	8	9					
0	0.018153	0.009070	0.003331					
1	0.010591	0.005048	0.002794					
2	0.008809	0.005000	0.002397					
3	0.003562	0.002545	0.001016					
0	0.000448	0.000461	0.000356					
1	0.000127	0.000087	0.000384					
2	0.000178	0.000513	0.000628					
3	0.000009	0.000090	0.000178					

Prediction of 6-digits occupation code (top 3)

```
[ ]: pred_top3_bottom = trans_bottom.predict(
    example_offer,
    output_level="last",
    top_k=3,
    format="dataframe"
)

pred_top3_top = trans_top.predict(
    example_offer,
    output_level="last",
    top_k=3,
    format="dataframe"
)
```

Initializing TransformerDataModule with model_name=FacebookAI/xlm-roberta-base,
max_seq_length=512, train/eval_batch_size=8/8, num_workers=2 ...

Setting up TransformerDataModule ...

/content/classifier/lib/python3.11/site-

packages/lightning_fabric/connector.py:572: `precision=16` is supported for
historical reasons but its usage is discouraged. Please set your precision to
16-mixed instead!

INFO:pytorch_lightning.utilities.rank_zero:Using 16bit Automatic Mixed Precision

```

(AMP)
INFO:pytorch_lightning.utilities.rank_zero:GPU available: True (cuda), used:
True
INFO:pytorch_lightning.utilities.rank_zero:TPU available: False, using: 0 TPU
cores
INFO:pytorch_lightning.utilities.rank_zero:HPU available: False, using: 0 HPUs
INFO:pytorch_lightning.utilities.rank_zero:Restoring states from the checkpoint
path at transformer-multi-
bottom-2024-sample-20/ckpts/epoch=6-val_loss=1.59784.ckpt

Starting predicting with TransformerClassifier ...

INFO:pytorch_lightning.utilities.migration.utils:Lightning automatically
upgraded your loaded checkpoint from v1.9.5 to v2.5.0.post0. To apply the
upgrade to your files permanently, run `python -m
pytorch_lightning.utilities.upgrade_checkpoint transformer-multi-
bottom-2024-sample-20/ckpts/epoch=6-val_loss=1.59784.ckpt`
/content/classifier/lib/python3.11/site-
packages/pytorch_lightning/callbacks/model_checkpoint.py:362: The dirpath has
changed from 'models-paper/multi/transformer-multi-bottom-2024-sample-20/ckpts'
to '/content/transformer-multi-
bottom-2024-sample-20/ckpts/epoch=6-val_loss=1.59784.ckpt', therefore
`best_model_score`, `kth_best_model_path`, `kth_value`, `last_model_path` and
`best_k_models` won't be reloaded. Only `best_model_path` will be reloaded.
INFO:pytorch_lightning.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES:
[0]
INFO:pytorch_lightning.utilities.rank_zero:Loaded model weights from the
checkpoint at transformer-multi-
bottom-2024-sample-20/ckpts/epoch=6-val_loss=1.59784.ckpt

Predicting: |          | 0/? [00:00<?, ?it/s]

Initializing TransformerDataModule with model_name=FacebookAI/xlm-roberta-base,
max_seq_length=512, train/eval_batch_size=8/8, num_workers=2 ...
Setting up TransformerDataModule ...

INFO:pytorch_lightning.utilities.rank_zero:Using 16bit Automatic Mixed Precision
(AMP)
INFO:pytorch_lightning.utilities.rank_zero:GPU available: True (cuda), used:
True
INFO:pytorch_lightning.utilities.rank_zero:TPU available: False, using: 0 TPU
cores
INFO:pytorch_lightning.utilities.rank_zero:HPU available: False, using: 0 HPUs
INFO:pytorch_lightning.utilities.rank_zero:Restoring states from the checkpoint
path at transformer-multi-top-2024-sample-20/ckpts/epoch=4-val_loss=0.31984.ckpt

Starting predicting with TransformerClassifier ...

INFO:pytorch_lightning.utilities.migration.utils:Lightning automatically
upgraded your loaded checkpoint from v1.9.5 to v2.5.0.post0. To apply the
upgrade to your files permanently, run `python -m

```

```

pytorch_lightning.utilities.upgrade_checkpoint transformer-multi-
top-2024-sample-20/ckpts/epoch=4-val_loss=0.31984.ckpt`
/content/classifier/lib/python3.11/site-
packages/pytorch_lightning/callbacks/model_checkpoint.py:362: The dirpath has
changed from 'models-paper/multi/transformer-multi-top-2024-sample-20/ckpts' to
'/content/transformer-multi-
top-2024-sample-20/ckpts/epoch=4-val_loss=0.31984.ckpt', therefore
`best_model_score`, `kth_best_model_path`, `kth_value`, `last_model_path` and
`best_k_models` won't be reloaded. Only `best_model_path` will be reloaded.
INFO:pytorch_lightning.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES:
[0]
INFO:pytorch_lightning.utilities.rank_zero:Loaded model weights from the
checkpoint at transformer-multi-
top-2024-sample-20/ckpts/epoch=4-val_loss=0.31984.ckpt

Predicting: |          | 0/? [00:00<?, ?it/s]

Results for top 3 along with the probabilities

```

```
[ ]: pd.concat([pred_top3_bottom,pred_top3_top])
```

```
[ ]:
  class_1 class_2 class_3   prob_1   prob_2   prob_3
0  431201  331401  252102  0.288475  0.193673  0.170916
1  242290  441990  411090  0.114219  0.080363  0.079738
2  251908  252102  251201  0.768940  0.035406  0.031738
3  251908  252102  252990  0.574634  0.329984  0.011291
0  212004  216590  431201  0.404771  0.082207  0.058343
1  242112  411090  242102  0.273998  0.130173  0.077134
2  251908  251201  251202  0.386948  0.224004  0.161463
3  252102  251908  252990  0.492204  0.336939  0.033593

```